

Team Project AAI 501

April 13, 2024

0.1 Introduction

The purpose of this project is to apply a classification algorithm to predict if someone will be classified as having a good or bad credit. The dataset consists of personal information such as age, sex, marital status, employment, credit amount request (among other) and classifies them based on these attributes as “good” and “bad” credit risks.

0.1.1 Business related questions

- What are some of the features that people with “Good” credit have?
- What is the average age of applicants?
- What is the average amount people have in their checking and savings account?
- What is the most common property type?
- what is the most common sex/status with “Good” credit
- Does an applicant’s demographics cause a deviation in the predicted category vs. what category they were actually placed in thus suggesting bias?

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from imblearn.over_sampling import SMOTE
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
from sklearn.feature_selection import SelectFromModel
```

```
[2]: df = pd.read_csv('german.data', sep='\s+')
df.head(2)
```

```
[2]:      A11    6  A34  A43  1169  A65  A75  4  A93  A101  ...  A121  67  A143  A152  \
0  A12  48  A32  A43  5951  A61  A73  2  A92  A101  ...  A121  22  A143  A152
1  A14  12  A34  A46  2096  A61  A74  2  A93  A101  ...  A121  49  A143  A152

      2  A173  1  A192  A201  1.1
0  1  A173  1  A191  A201    2
1  1  A172  2  A191  A201    1

[2 rows x 21 columns]
```

```
[3]: column_names = ['checking_account', 'duration_month', 'credit_history',
    ↪ 'credit_purpose', 'credit_amount', 'savings_account', 'present_employment', 'disposable_income',
    ↪
    ↪ 'status_sex', 'debtors', 'residence_since', 'property', 'age', 'other_installments',
    ↪ 'housing', 'credits_at_current_bank', 'job', 'dependants', 'telephone',
    ↪ 'foreign_worker', 'class']
```

```
[4]: df.columns = column_names
df.head(2)
```

```
[4]:  checking_account  duration_month  credit_history  credit_purpose  \
0              A12              48              A32              A43
1              A14              12              A34              A46

      credit_amount  savings_account  present_employment  \
0              5951              A61              A73
1              2096              A61              A74

disposable_income_percent  status_sex  debtors  ...  property  age  \
0              2              A92  A101  ...  A121  22
1              2              A93  A101  ...  A121  49

      other_installments  housing  credits_at_current_bank  job  dependants  \
0              A143  A152              1  A173              1
1              A143  A152              1  A172              2

      telephone  foreign_worker  class
0      A191              A201    2
1      A191              A201    1

[2 rows x 21 columns]
```

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 999 entries, 0 to 998
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
#   ...  ...
```

```

---  -----
0  checking_account      999 non-null  object
1  duration_month        999 non-null  int64
2  credit_history         999 non-null  object
3  credit_purpose           999 non-null  object
4  credit_amount          999 non-null  int64
5  savings_account        999 non-null  object
6  present_employment     999 non-null  object
7  disposable_income_percent 999 non-null  int64
8  status_sex             999 non-null  object
9  debtors                999 non-null  object
10 residence_since        999 non-null  int64
11 property               999 non-null  object
12 age                    999 non-null  int64
13 other_installments     999 non-null  object
14 housing                 999 non-null  object
15 credits_at_current_bank 999 non-null  int64
16 job                    999 non-null  object
17 dependants             999 non-null  int64
18 telephone              999 non-null  object
19 foreign_worker          999 non-null  object
20 class                  999 non-null  int64

```

dtypes: int64(8), object(13)

memory usage: 164.0+ KB

```
[6]: df.describe()
```

```

[6]:      duration_month  credit_amount  disposable_income_percent  \
count      999.000000      999.000000      999.000000
mean       20.917918      3273.362362       2.971972
std        12.055619      2823.365811       1.118802
min         4.000000       250.000000       1.000000
25%        12.000000      1368.500000       2.000000
50%        18.000000      2320.000000       3.000000
75%        24.000000      3972.500000       4.000000
max        72.000000     18424.000000       4.000000

      residence_since      age  credits_at_current_bank  dependants  \
count      999.000000      999.000000      999.000000      999.000000
mean         2.843844      35.514515         1.406406       1.155155
std          1.103665      11.337487         0.577639       0.362234
min           1.000000      19.000000         1.000000       1.000000
25%           2.000000      27.000000         1.000000       1.000000
50%           3.000000      33.000000         1.000000       1.000000
75%           4.000000      42.000000         2.000000       1.000000
max           4.000000      75.000000         4.000000       2.000000

```

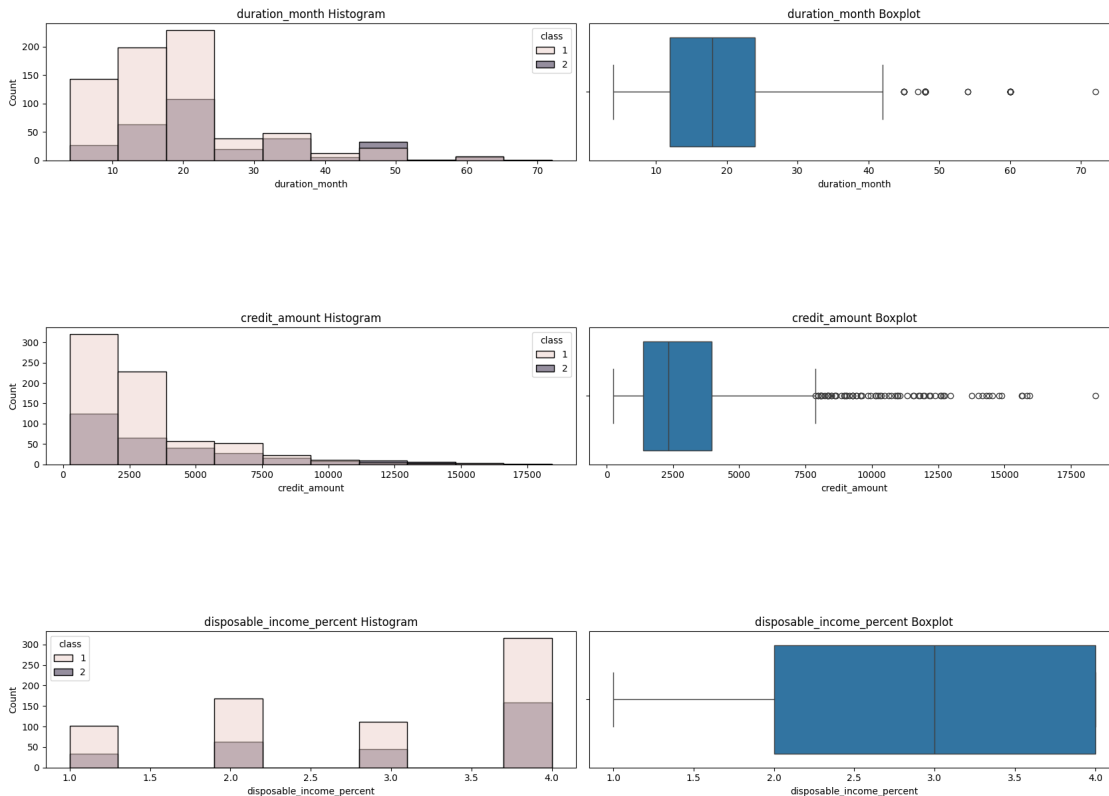
	class
count	999.000000
mean	1.300300
std	0.458618
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	2.000000

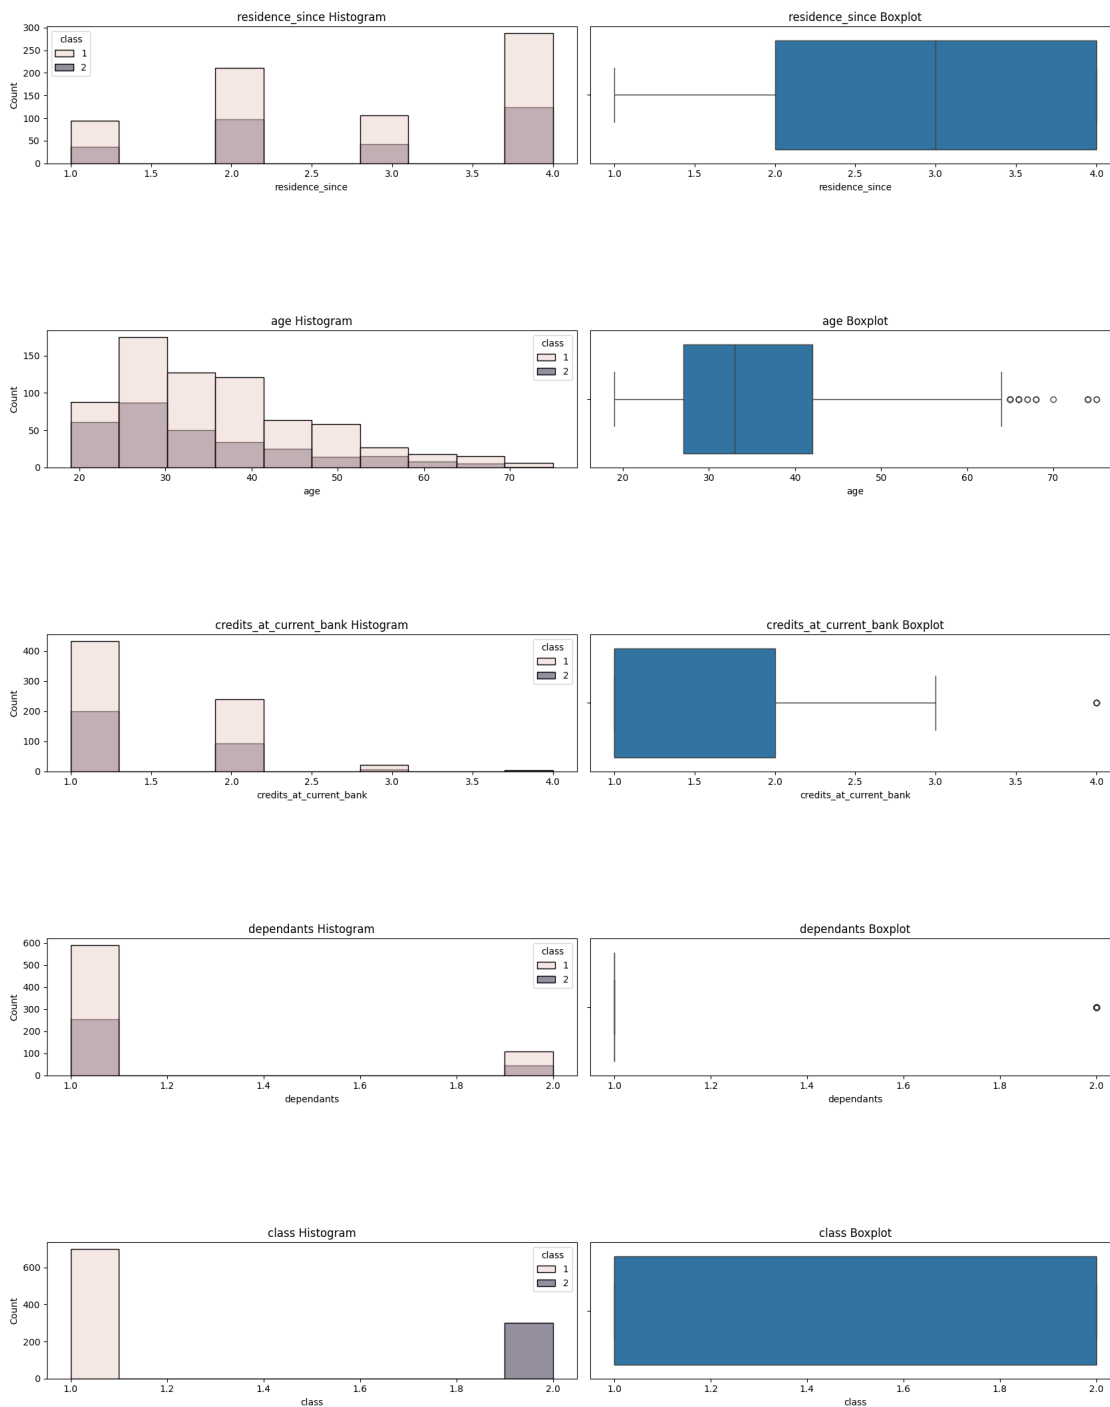
```
[7]: for i in df.columns:
      if df[i].dtype == 'int64':
          fig, ax = plt.subplots(1, 2, figsize=(17, 3))

          sns.histplot(data=df, x=i, bins=10, ax=ax[0], hue='class')
          ax[0].set_title(f'{i} Histogram')

          sns.boxplot(data=df, x=i, ax=ax[1])
          ax[1].set_title(f'{i} Boxplot')

          plt.tight_layout()
          plt.show();
```





```
[8]: object_columns = df.select_dtypes(include='object').columns
```

```
num_rows = (len(object_columns) + 1) // 2
```

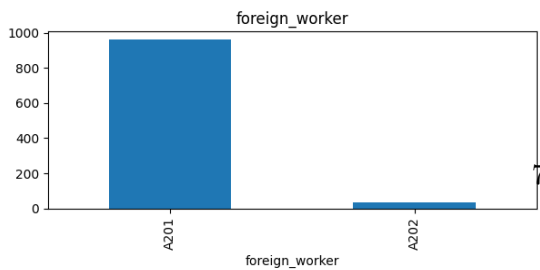
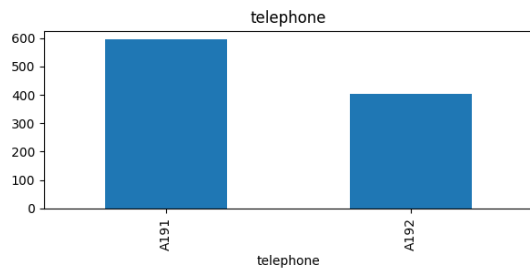
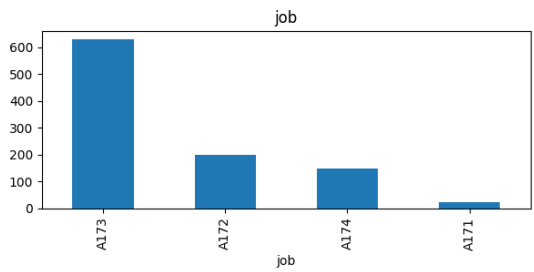
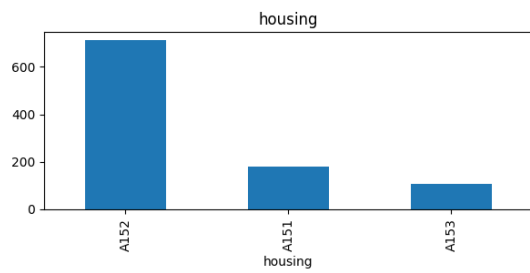
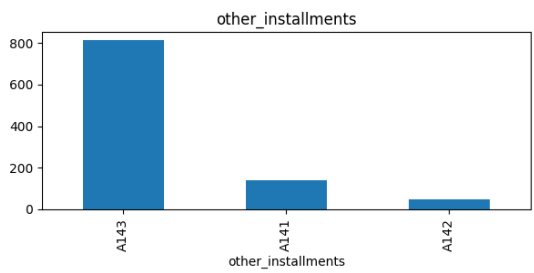
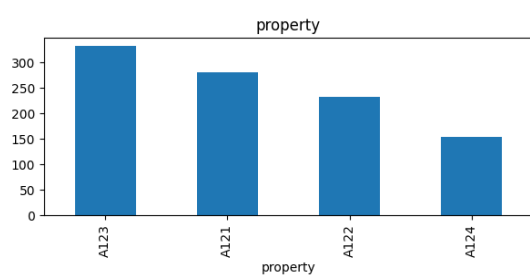
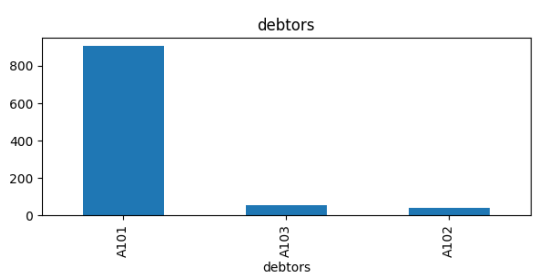
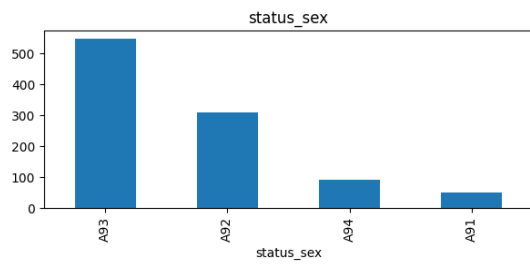
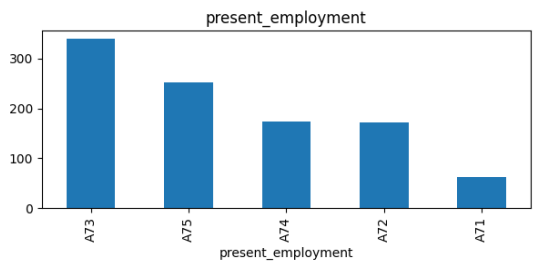
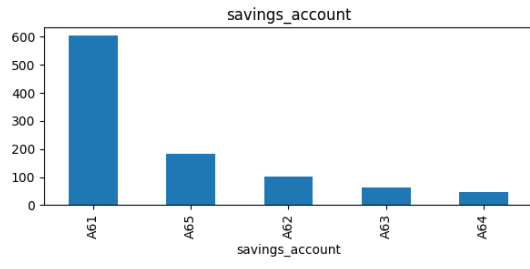
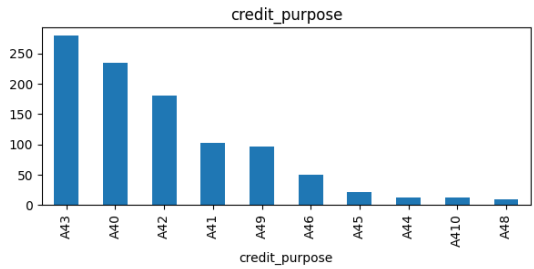
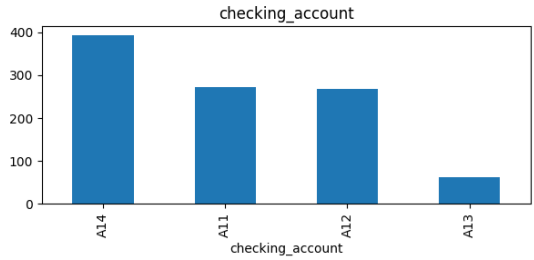
```
fig, axes = plt.subplots(num_rows, 2, figsize=(12, 3*num_rows))

axes = axes.flatten()

for i, col in enumerate(object_columns):
    df[col].value_counts().plot(kind='bar', ax=axes[i])
    axes[i].set_title(col)

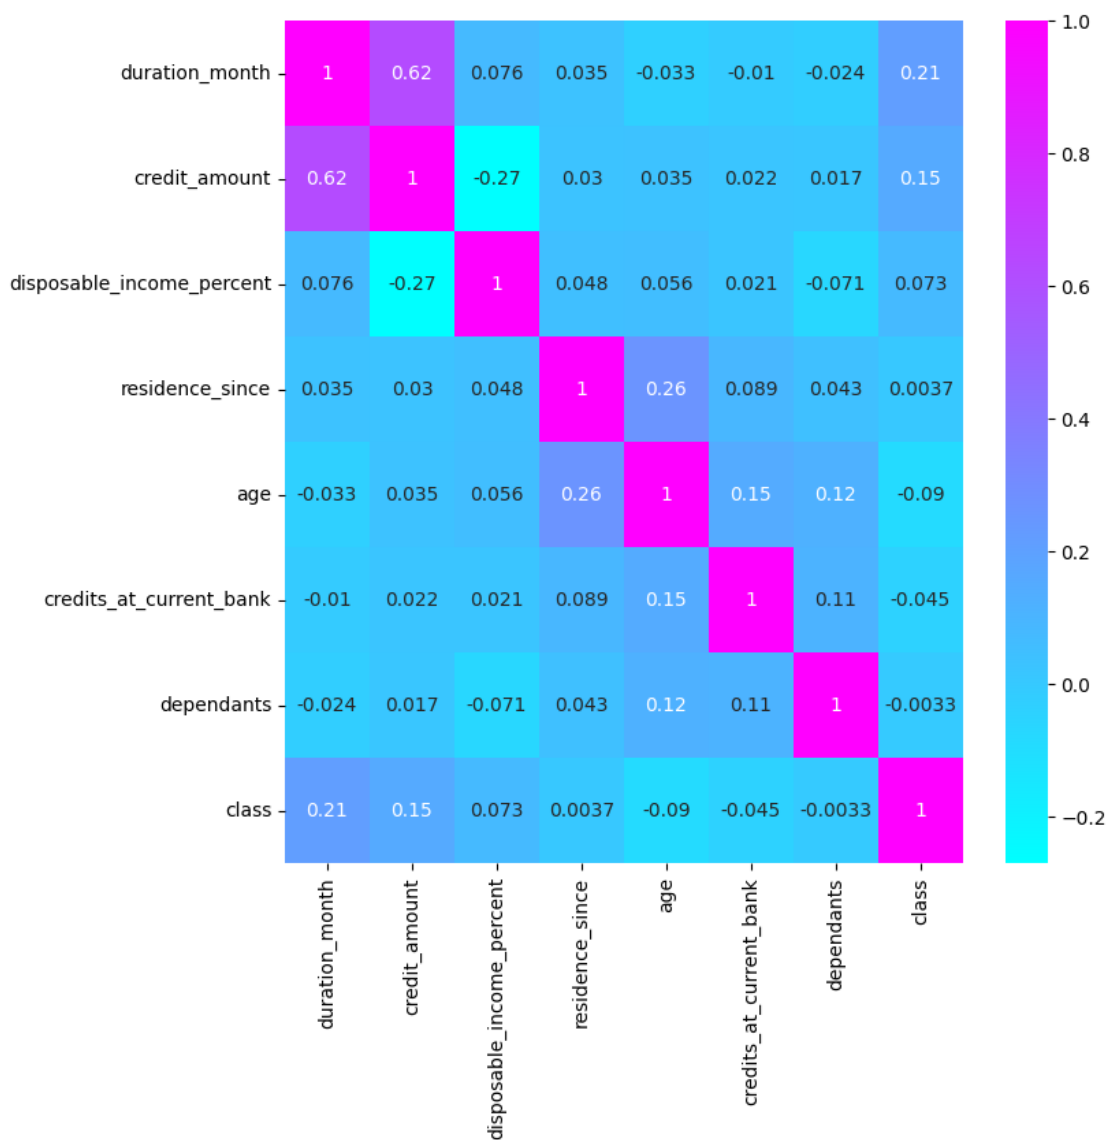
for i in range(len(object_columns), len(axes)):
    axes[i].axis('off')

plt.tight_layout()
plt.show()
```



```
[9]: corr_matr = df.corr(numeric_only=True)
plt.figure(figsize=(8,8))
sns.heatmap(corr_matr, cmap= 'cool', annot= True)
```

[9]: <Axes: >



0.1.2 Analyzing Good vs Bad Credit

```
[10]: bad_credit = df[df['class'] == 2]
      good_credit = df[df['class'] == 1]
```

```
[11]: bad_credit.groupby('credits_at_current_bank').mean(numeric_only = True)
```

```
[11]:
```

	duration_month	credit_amount	\
credits_at_current_bank			
1	24.335000	3751.870000	
2	26.076087	4313.076087	
3	24.000000	4204.000000	
4	24.000000	4518.500000	

	disposable_income_percent	residence_since	\
credits_at_current_bank			
1	3.150000	2.770000	
2	3.032609	2.956522	
3	2.333333	3.500000	
4	3.000000	4.000000	

	age	dependants	class
credits_at_current_bank			
1	33.250000	1.135000	2.0
2	35.054348	1.184783	2.0
3	40.833333	1.333333	2.0
4	34.500000	1.000000	2.0

```
[12]: good_credit.groupby('credits_at_current_bank').mean(numeric_only = True)
```

```
[12]:
```

	duration_month	credit_amount	\
credits_at_current_bank			
1	19.166282	2963.034642	
2	19.858333	3043.820833	
3	13.454545	3008.772727	
4	19.500000	2236.750000	

	disposable_income_percent	residence_since	\
credits_at_current_bank			
1	2.856813	2.806005	
2	3.029167	2.850000	
3	2.727273	3.272727	
4	4.000000	3.750000	

	age	dependants	class
credits_at_current_bank			
1	35.247113	1.129330	1.0
2	36.716667	1.179167	1.0

3	45.545455	1.363636	1.0
4	53.500000	1.500000	1.0

Observations: - The **average credit amount** in people with bad credit tends to be more in general in comparison to people with good credit. In other words, bad credit has more accounts open and ask for more money - The **average age** of people with 4 or more credit accounts at the current bank is 53 years old with good credit and 34.5 years old with bad credit - We can drop the dependants and residence_since. These features are very similar between people with good and bad credit.

```
[13]: df = df.drop(columns=['dependants', 'residence_since'], axis= 1)
```

Observation:

I am dropping Dependats and residence since there is no strong correlation among other variables.

Savings account/bonds

A61 : ... < 100 DM

A62 : 100 <= ... < 500 DM

A63 : 500 <= ... < 1000 DM

A64 : .. >= 1000 DM

A65 : unknown/ no savings account

Status of existing checking account

A11 : ... < 0 DM

A12 : 0 <= ... < 200 DM

A13 : ... >= 200 DM /

salary assignments for at least 1 year

A14 : no checking account

```
[14]: print(f'Good credit:\n{good_credit.savings_account.value_counts()}')
print(f'Bad Credit:\n{bad_credit.savings_account.value_counts()}')
print(f'Good credit:\n{good_credit.checking_account.value_counts()}')
print(f'Bad Credit:\n{bad_credit.checking_account.value_counts()}')
```

Good credit:

savings_account

A61 386

A65 150

A62 69

A63 52

A64 42

Name: count, dtype: int64

Bad Credit:

savings_account

```

A61    217
A62     34
A65     32
A63     11
A64      6
Name: count, dtype: int64
Good credit:
checking_account
A14     348
A12     164
A11     138
A13      49
Name: count, dtype: int64
Bad Credit:
checking_account
A11     135
A12     105
A14      46
A13      14
Name: count, dtype: int64

```

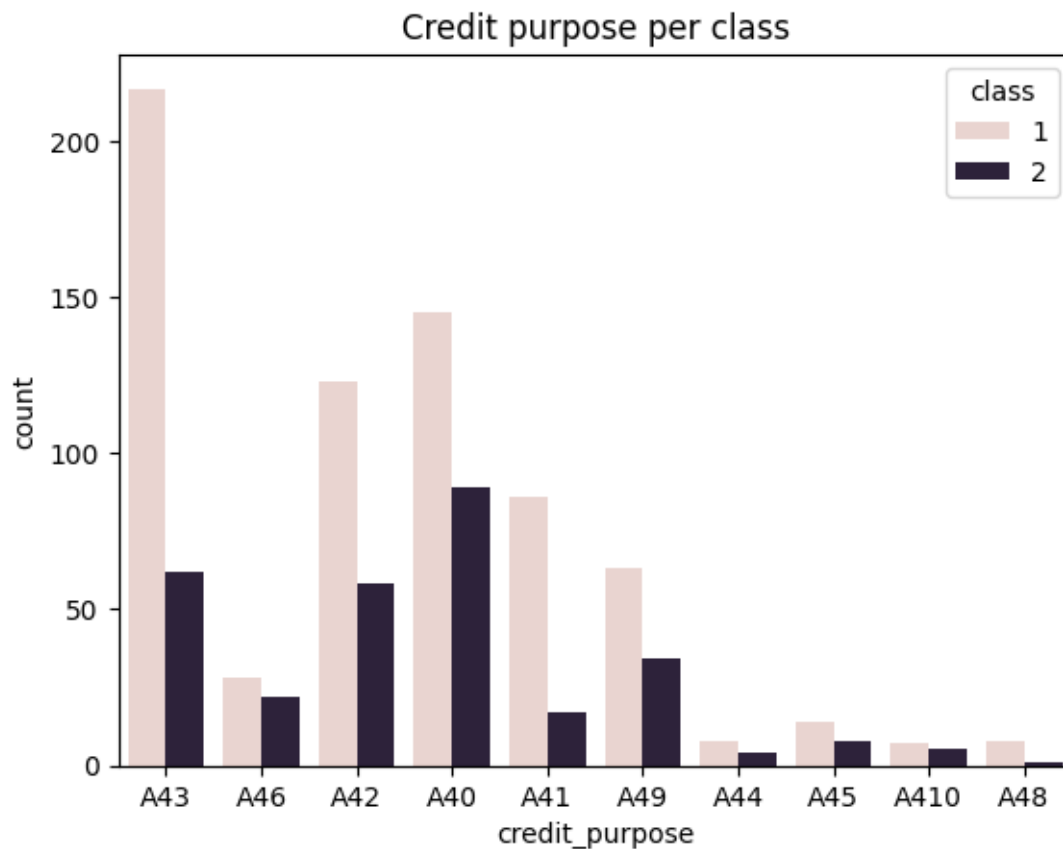
Observations: - The majority of people with good credit do not have checking account at the current bank, while people with bad credit have more accoutns open but have less than 200 DM - For the people that do have an account open, people with good credit have more than 200 DM in their account

Maybe We should consider using one-hot encoding on checking and savings account

```

[15]: sns.countplot(x='credit_purpose',hue='class', data= df)
      plt.title('Credit purpose per class')
      plt.show();

```



Attribute 4: (qualitative)

Purpose

A40 : car (new)

A41 : car (used)

A42 : furniture/equipment

A43 : radio/television

A44 : domestic appliances

A45 : repairs

A46 : education

A47 : (vacation - does not exist?)

A48 : retraining

A49 : business

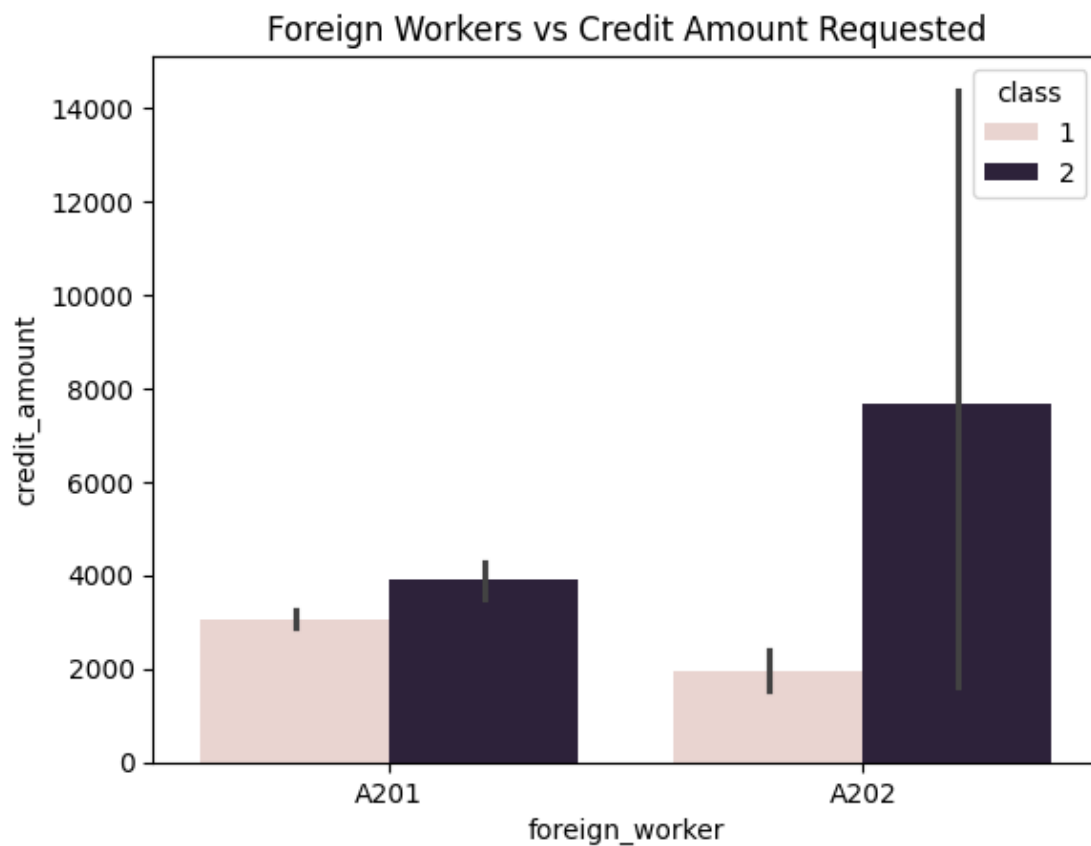
A410 : others

Observations: - Class 1(good) use their credit for radio and televisions, cars and furniture/equipment, while people with bad credit tend to use their credit mostly for purchasing a car

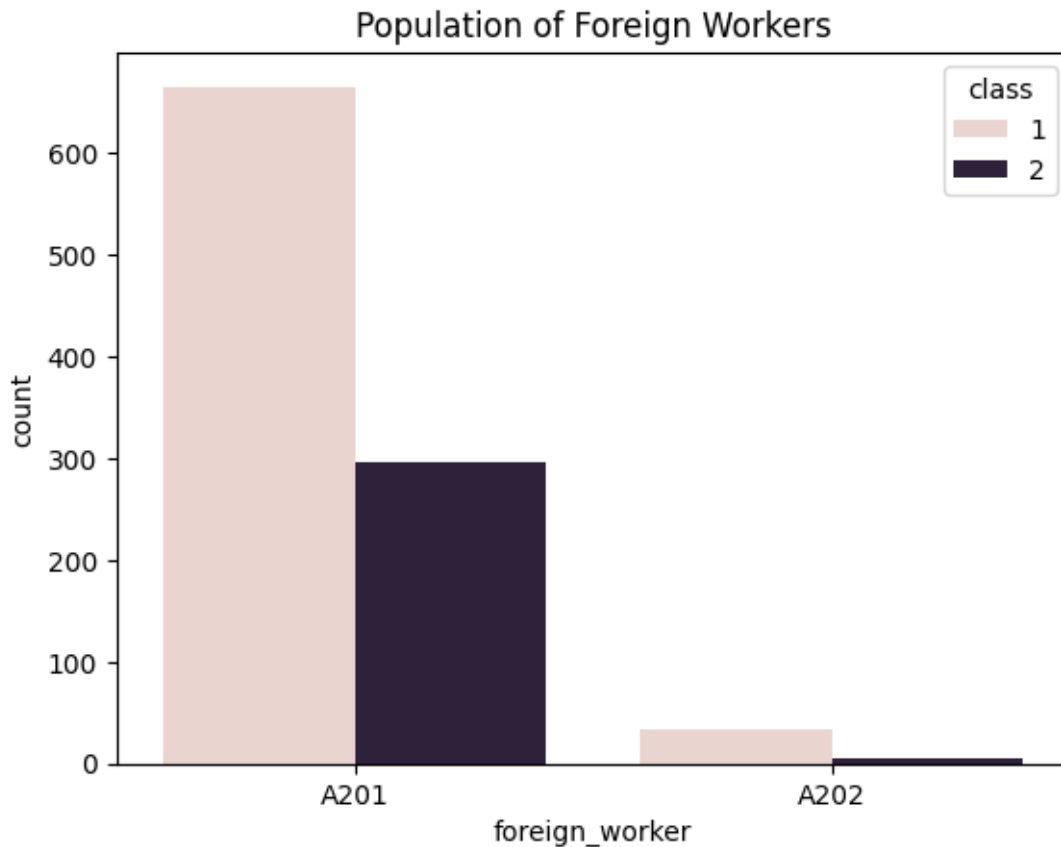
We should use one-hot encoding for credit purpose

```
[16]: df_encoded = pd.get_dummies(df, columns=['credit_purpose'], prefix='credit_purpose')
      ↪ 'credit_purpose')

[17]: sns.barplot(x = 'foreign_worker', y = 'credit_amount', data = df, hue='class')
      plt.title('Foreign Workers vs Credit Amount Requested')
      plt.show();
```



```
[18]: sns.countplot(x=df['foreign_worker'], hue = df['class'])
      plt.title('Population of Foreign Workers')
      plt.show();
```



```
[19]: print(df['foreign_worker'].value_counts().sum())
      print('-----')
      print(df['foreign_worker'].value_counts()/999)
```

999

foreign_worker

A201 0.962963

A202 0.037037

Name: count, dtype: float64

foreign worker:

A201 : yes

A202 : no

Observations: - Foreign workers, tend to ask higher ammounts of credit and tend to have a population with higher bad credit history

- Since Foreign workers tend to have bad credit we should encode this as well and becasue it composes 96% of our dataset

```
[20]: # Dropping Foreign Column to Avoid Multicollinearity
df_encoded = pd.get_dummies(df_encoded, columns=['foreign_worker'],
                             drop_first=True, prefix='Foreign')
```

```
[21]: sns.countplot(x = 'credit_history', data=df, hue = 'class')
plt.title('Credit History')
plt.show();
```



Credit history

A30 : no credits taken/

all credits paid back duly

A31 : all credits at this bank paid back duly

A32 : existing credits paid back duly till now

A33 : delay in paying off in the past

A34 : critical account/ other credits existing (not at this bank)

Observations: - Although the majority of people with good and bad credit tend to pay off all their debt. There is a significant amount of people with bad debt that have accounts classified as

critical in other banks. This could mean that they are looking to get more credit with bad history in other banks

```
[22]: df_encoded = pd.get_dummies(df_encoded, columns=['credit_history'],  
    ↪ prefix='credit_history')
```

```
[23]: sns.countplot(data=df, x='debtors', hue='class')  
plt.title('Guarantors/Debtors')  
plt.show()
```



Other debtors / guarantors

A101 : none

A102 : co-applicant

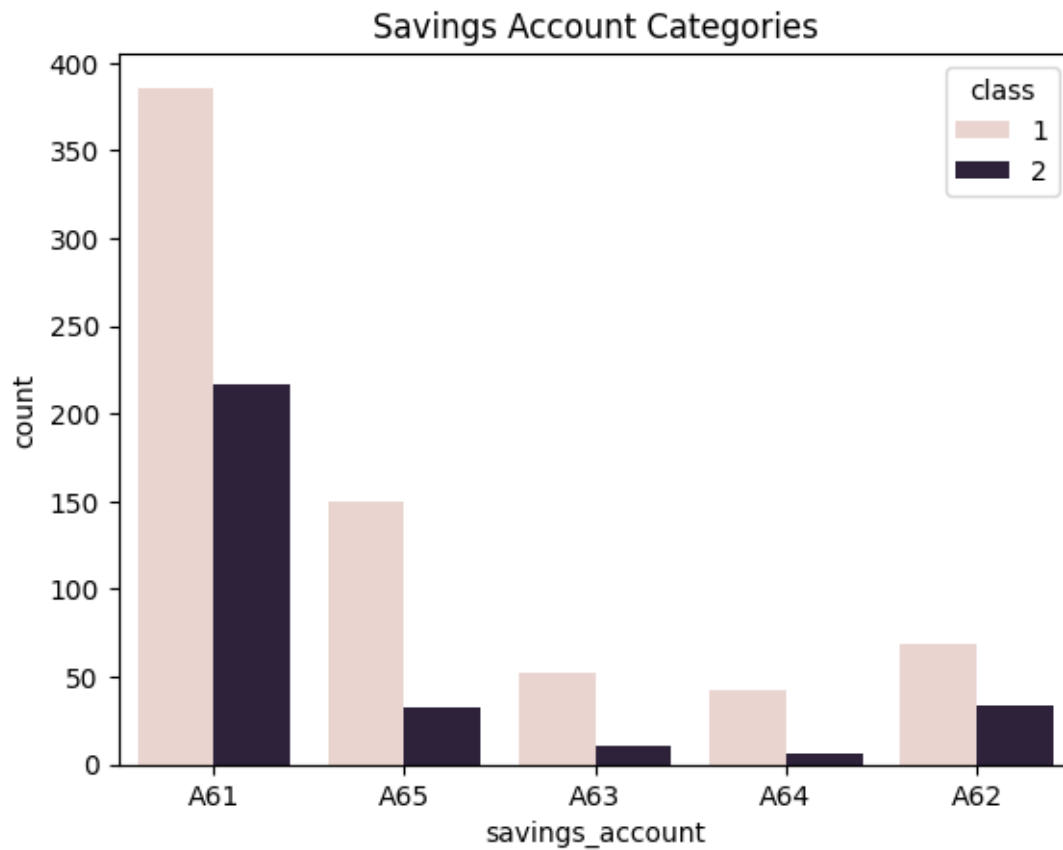
A103 : guarantor

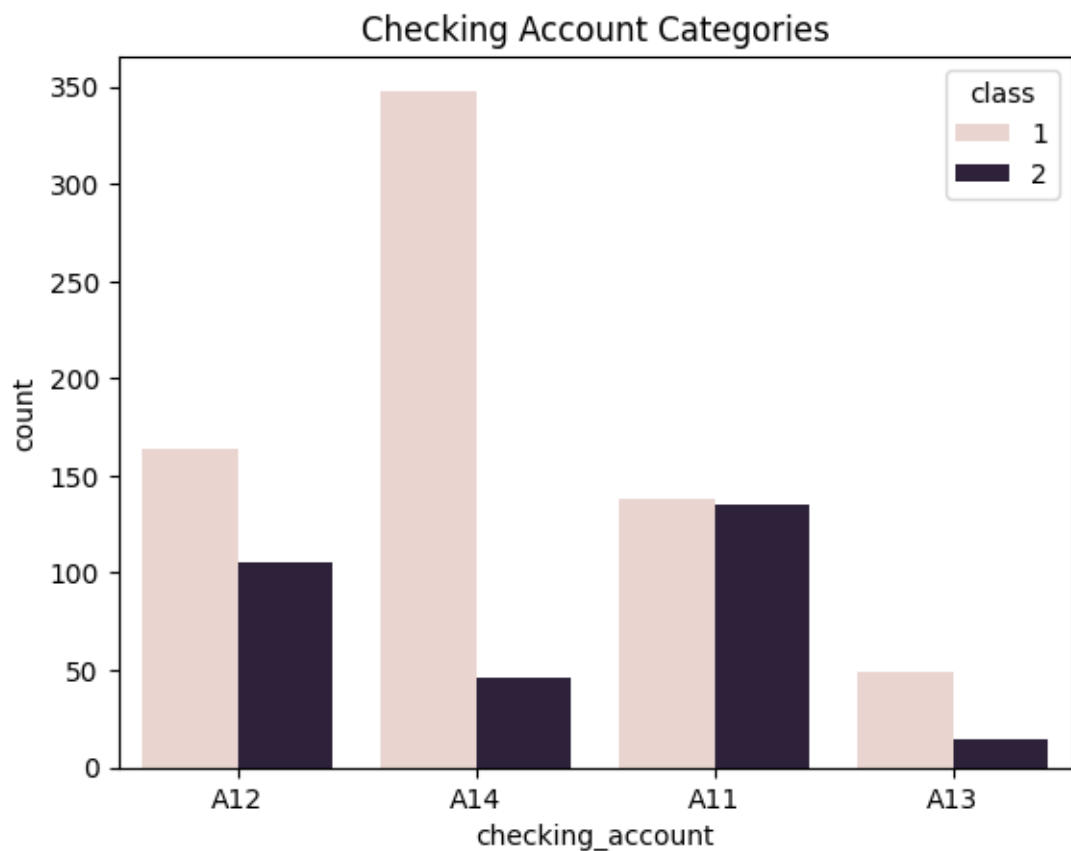
We can eliminate this column

```
[24]: df_encoded = df_encoded.drop(columns= 'debtors', axis=1)
```



```
[25]: sns.countplot(data = df , x = df['savings_account'], hue='class')
plt.title('Savings Account Categories')
plt.show();
sns.countplot(data = df , x = df['checking_account'], hue='class')
plt.title('Checking Account Categories')
plt.show();
```





checking account

A11 : ... < 0 DM

A12 : 0 <= ... < 200 DM

A13 : ... >= 200 DM /
salary assignments for at least 1 year

A14 : no checking account

Savings account/bonds

A61 : ... < 100 DM

A62 : 100 <= ... < 500 DM

A63 : 500 <= ... < 1000 DM

A64 : .. >= 1000 DM

A65 : unknown/ no savings account

We can include checking and savings account feature since it is important for financial institutions to evaluate how much available capital they have

```
[26]: df_encoded = pd.get_dummies(df_encoded, columns=['checking_account'],  
    ↪ prefix='checking_account')  
df_encoded = pd.get_dummies(df_encoded, columns=['savings_account'],  
    ↪ prefix='savings_account')
```

Dropping telephone since it is not a determinant factor for credit risk worthiness

```
[27]: df_encoded = df_encoded.drop(columns='telephone' , axis= 1)
```

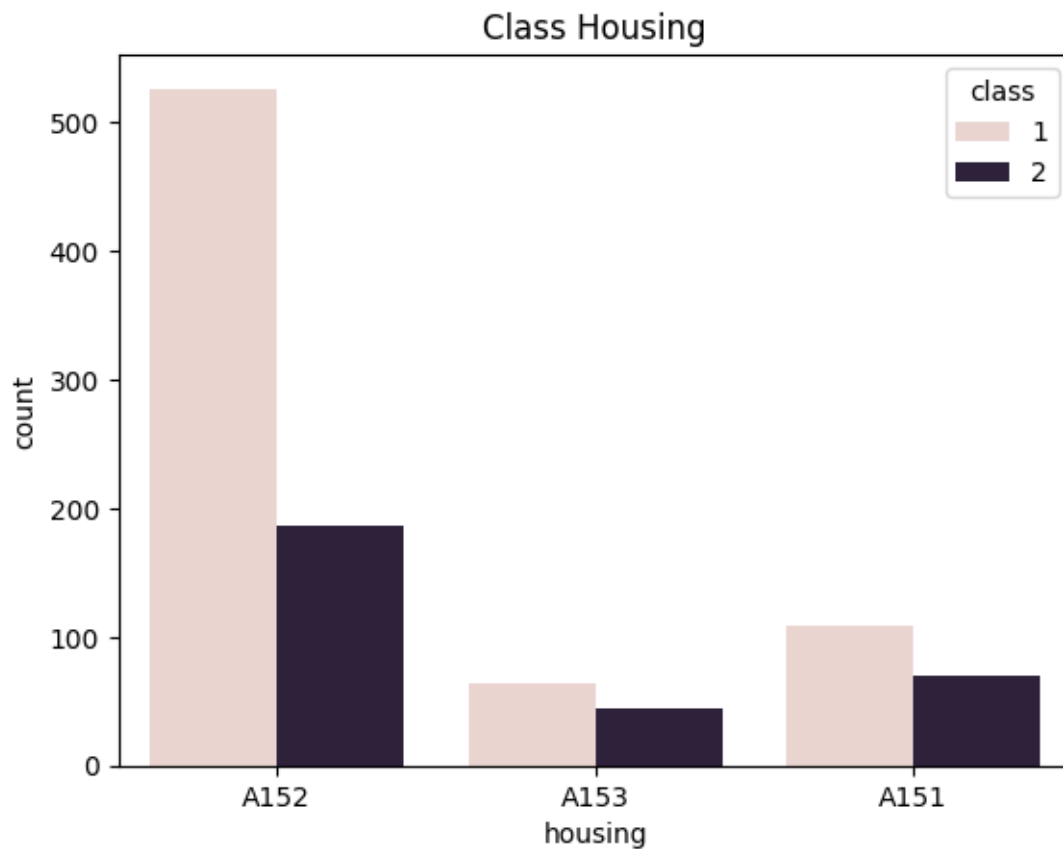
```
[28]: df.groupby('credits_at_current_bank')['class'].sum()
```

```
[28]: credits_at_current_bank  
1      833  
2      424  
3       34  
4        8  
Name: class, dtype: int64
```

Double of the amount of people with good credit have only one account open at the current bank in comparison to people with bad credit

```
[29]: df_encoded = pd.get_dummies(df_encoded, columns=['credits_at_current_bank'],  
    ↪ prefix='current_bank_credit')
```

```
[30]: sns.countplot(data=df_encoded, x = 'housing', hue='class')  
plt.title('Class Housing')  
plt.show();
```

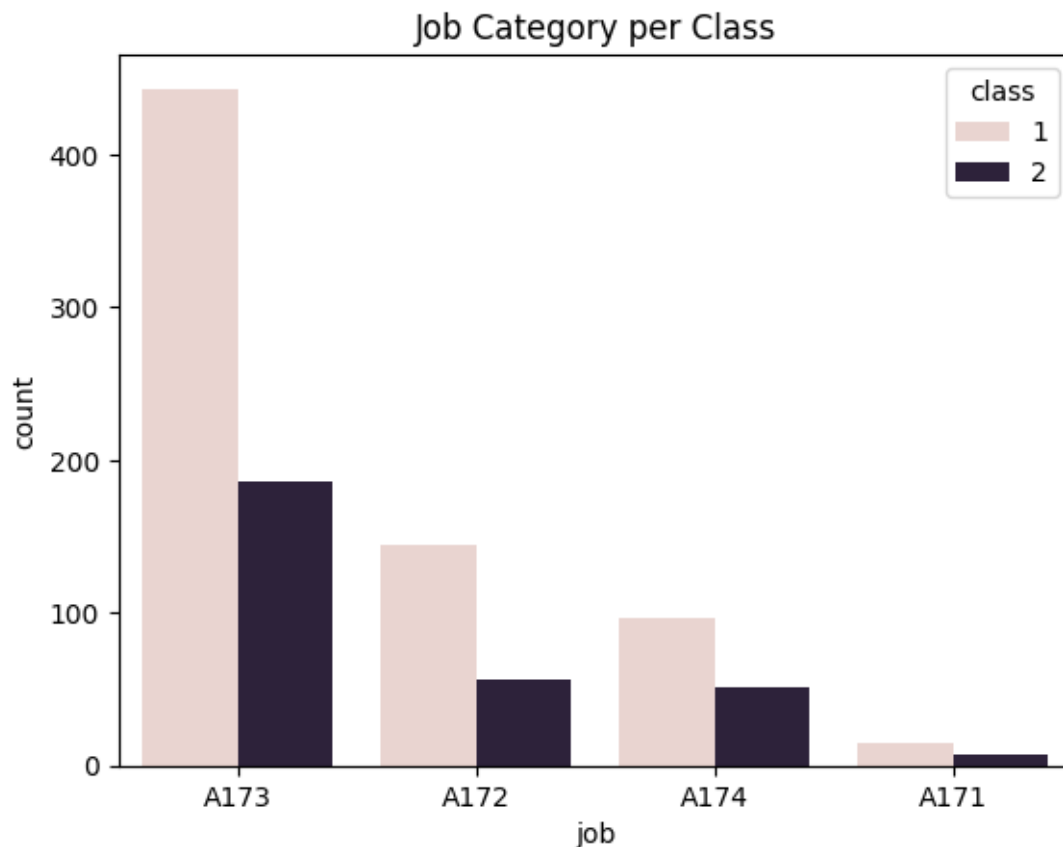


Housing
 A151 : rent
 A152 : own
 A153 : for free

Owning a home is an important factor for credit approval and usually homeowners need to have a 'decent' credit to own a home

```
[31]: df_encoded = pd.get_dummies(df_encoded, columns=['housing'], prefix='housing')
```

```
[32]: sns.countplot(data=df_encoded, x = 'job', hue = 'class' )
plt.title('Job Category per Class')
plt.show();
```

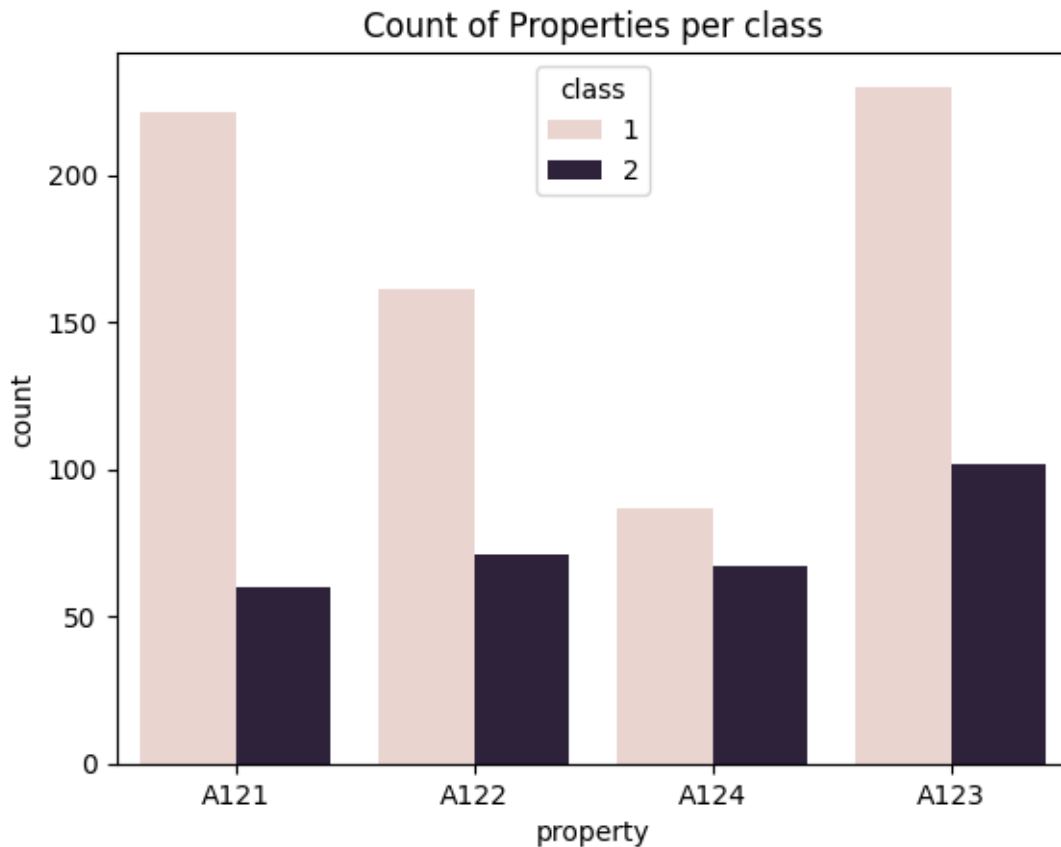


Job
 A171 : unemployed/ unskilled - non-resident
 A172 : unskilled - resident
 A173 : skilled employee / official
 A174 : management/ self-employed/
 highly qualified employee/ officer

The distribution for each class and employment seems proportional, we can omit this category

```
[33]: df_encoded = pd.get_dummies(df_encoded, columns=['job'], prefix='job')
```

```
[34]: sns.countplot(data=df_encoded, x = 'property', hue = 'class')
plt.title('Count of Properties per class')
plt.show();
```



Property

A121 : real estate

A122 : if not A121 : building society savings agreement/life insurance

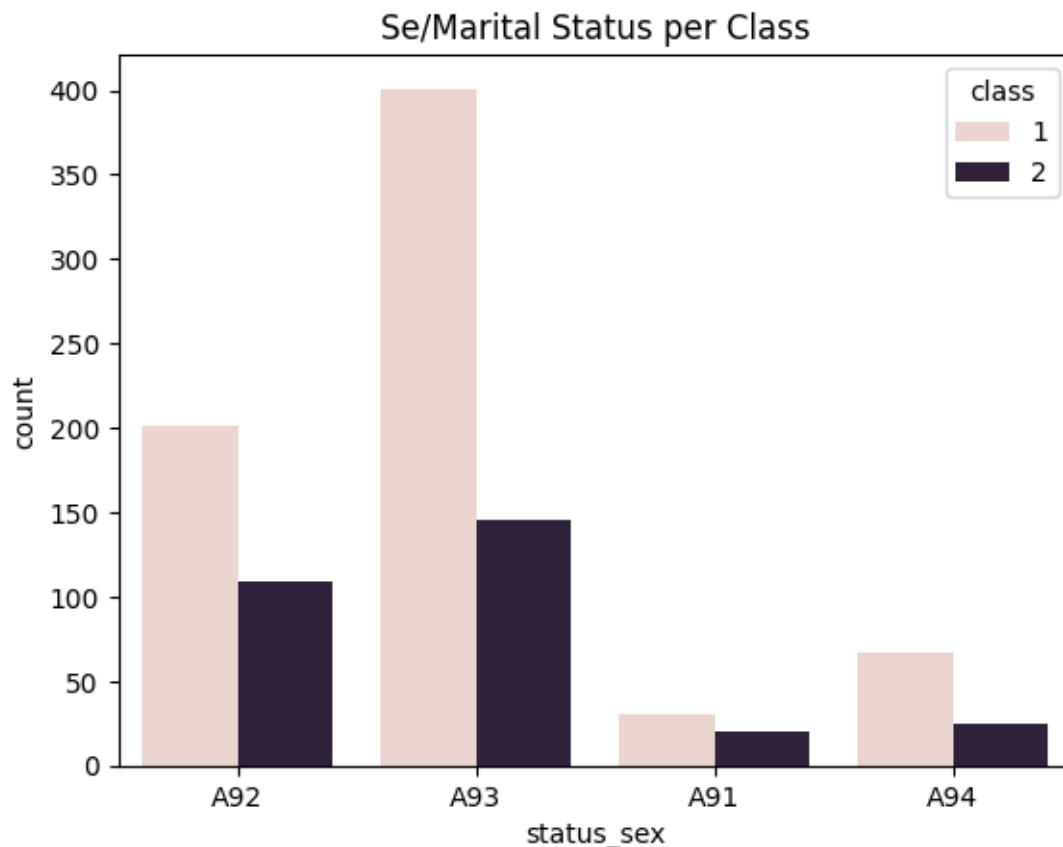
A123 : if not A121/A122 : car or other, not in attribute 6

A124 : unknown / no property

Encoding this feature because banks tend to look at collaterals when applying for credit lines

```
[35]: df_encoded = pd.get_dummies(df_encoded, columns=['property'], prefix='property')
```

```
[36]: sns.countplot(data=df_encoded, x = 'status_sex', hue = 'class')
plt.title('Se/Marital Status per Class')
plt.show();
```



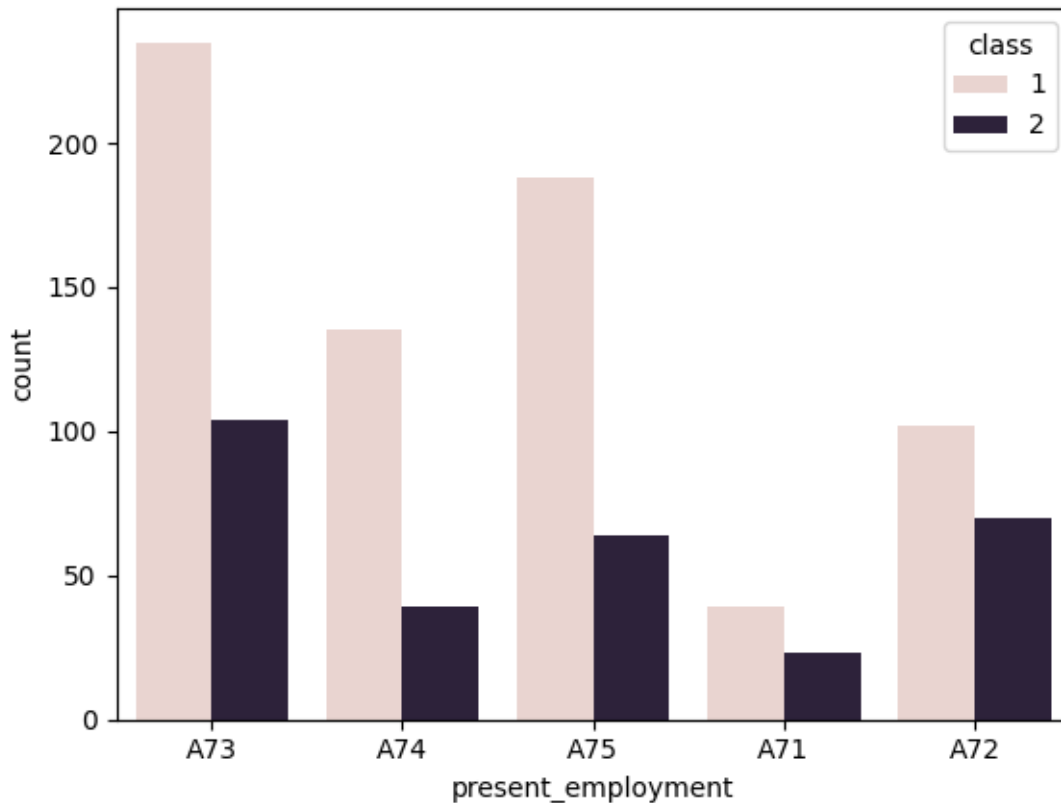
Personal status and sex
 A91 : male : divorced/separated
 A92 : female : divorced/separated/married
 A93 : male : single
 A94 : male : married/widowed
 A95 : female : single

Encoding this feature because you can tell there is a difference between good/bad credit between male and females

```
[37]: df_encoded = pd.get_dummies(df_encoded, columns=['status_sex'],
    ↳ prefix='status_sex')
```

```
[38]: sns.countplot(data=df_encoded, x = 'present_employment', hue = 'class')
```

```
[38]: <Axes: xlabel='present_employment', ylabel='count'>
```



```

A71 : unemployed
A72 :      ... < 1 year
A73 : 1  <= ... < 4 years
A74 : 4  <= ... < 7 years
A75 :      .. >= 7 years

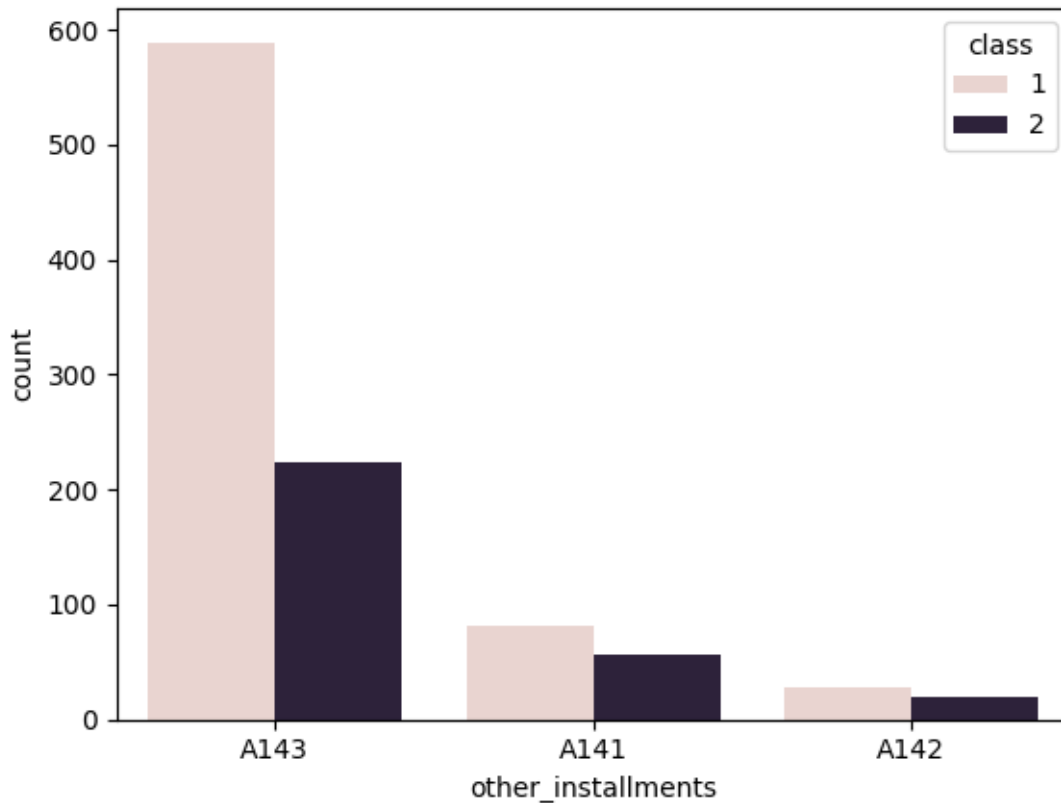
```

Encoding this feature because it is important to have a working history that is stable, it could determine your eligibility for a credit line (domain knowledge)

```
[39]: df_encoded = pd.get_dummies(df_encoded, columns=['present_employment'],
    ↪ prefix='years_in employment')
```

```
[40]: sns.countplot(data=df_encoded, x = 'other_installments', hue = 'class')
```

```
[40]: <Axes: xlabel='other_installments', ylabel='count'>
```

Other installment plans

A141 : bank
A142 : stores
A143 : none

Having Other Installments does not seem that affect wether you have good or bad credit,since the population results look proportional

```
[41]: #df_encoded = pd.get_dummies(df_encoded, columns=['other_installments'],
      ↪ prefix='other_installments')
df_encoded = df_encoded.drop(columns='other_installments', axis= 1)
```

```
[42]: pd.set_option('display.max_columns', None)

df_encoded.head(1)
```

```
[42]: duration_month  credit_amount  disposable_income_percent  age  class  \
0          48          5951                2  22      2

credit_purpose_A40  credit_purpose_A41  credit_purpose_A410  \
0          False          False          False
```

0	credit_purpose_A42	credit_purpose_A43	credit_purpose_A44	\		
	False	True	False			
0	credit_purpose_A45	credit_purpose_A46	credit_purpose_A48	\		
	False	False	False			
0	credit_purpose_A49	Foreign_A202	credit_history_A30	credit_history_A31 \		
	False	False	False	False		
0	credit_history_A32	credit_history_A33	credit_history_A34	\		
	True	False	False			
0	checking_account_A11	checking_account_A12	checking_account_A13	\		
	False	True	False			
0	checking_account_A14	savings_account_A61	savings_account_A62	\		
	False	True	False			
0	savings_account_A63	savings_account_A64	savings_account_A65	\		
	False	False	False			
0	current_bank_credit_1	current_bank_credit_2	current_bank_credit_3	\		
	True	False	False			
0	current_bank_credit_4	housing_A151	housing_A152	housing_A153 job_A171 \		
	False	False	True	False False		
0	job_A172	job_A173	job_A174	property_A121	property_A122	property_A123 \
	False	True	False	True	False	False
0	property_A124	status_sex_A91	status_sex_A92	status_sex_A93	\	
	False	False	True	False		
0	status_sex_A94	years_in employment_A71	years_in employment_A72	\		
	False	False	False			
0	years_in employment_A73	years_in employment_A74	years_in employment_A75			
	True	False	False			

0.2 Machine Learning

```
[43]: #1 is Good, 0 is bad
mapping_dict = {1: 1, 2: 0}
df_encoded['class'] = df_encoded['class'].map(mapping_dict)

#Seperating Features from Target
x = df_encoded.drop(columns='class',axis =1)
```

```
y = df_encoded['class']
```

0.3 Naive Bayes Algorithm With Standardization

```
[44]: #Splitting
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= .60,
    ↪random_state= 42)

#Standardizing
scaler = StandardScaler()
xtrain_scale = scaler.fit_transform(x_train)
xtest = scaler.transform(x_test)

# Hyperparameter Tuning
param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
    ↪cv=10, scoring='accuracy')
grid_search.fit(xtrain_scale, y_train)

# Training
n_b_best = grid_search.best_estimator_
n_b_best.fit(xtrain_scale, y_train)
predictions = n_b_best.predict(xtest)

#Performance
scores = n_b_best.score(xtest, y_test)
conf_mtrx = confusion_matrix(y_test, predictions)
precision = precision_score(y_test, predictions)
recall = recall_score(y_test, predictions)
f1 = f1_score(y_test, predictions)
disp = ConfusionMatrixDisplay(confusion_matrix=conf_mtrx)

#Plotting ROC Curve
y_pred_proba = n_b_best.predict_proba(xtest)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba, pos_label=1)
roc_auc = auc(fpr, tpr)

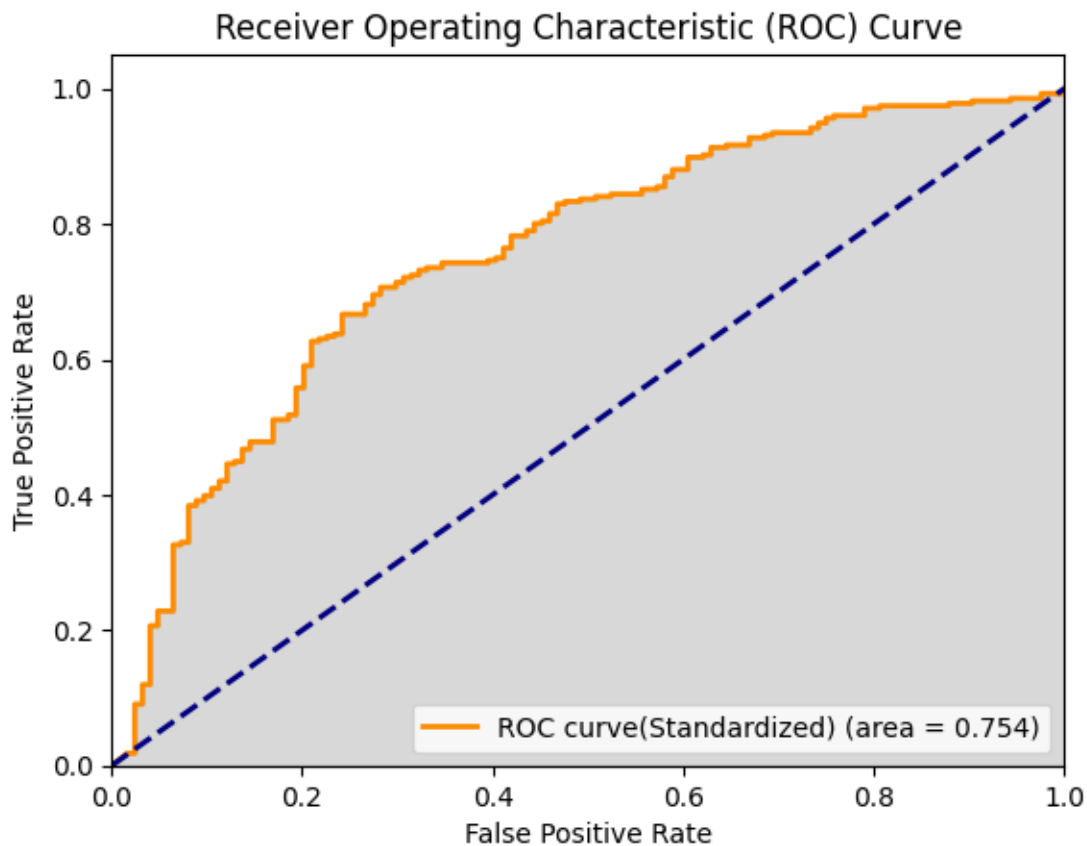
#Results
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve(Standardized)
    ↪(area = {round(roc_auc,3)})')
plt.fill_between(fpr, tpr, color='gray', alpha=0.3)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

```

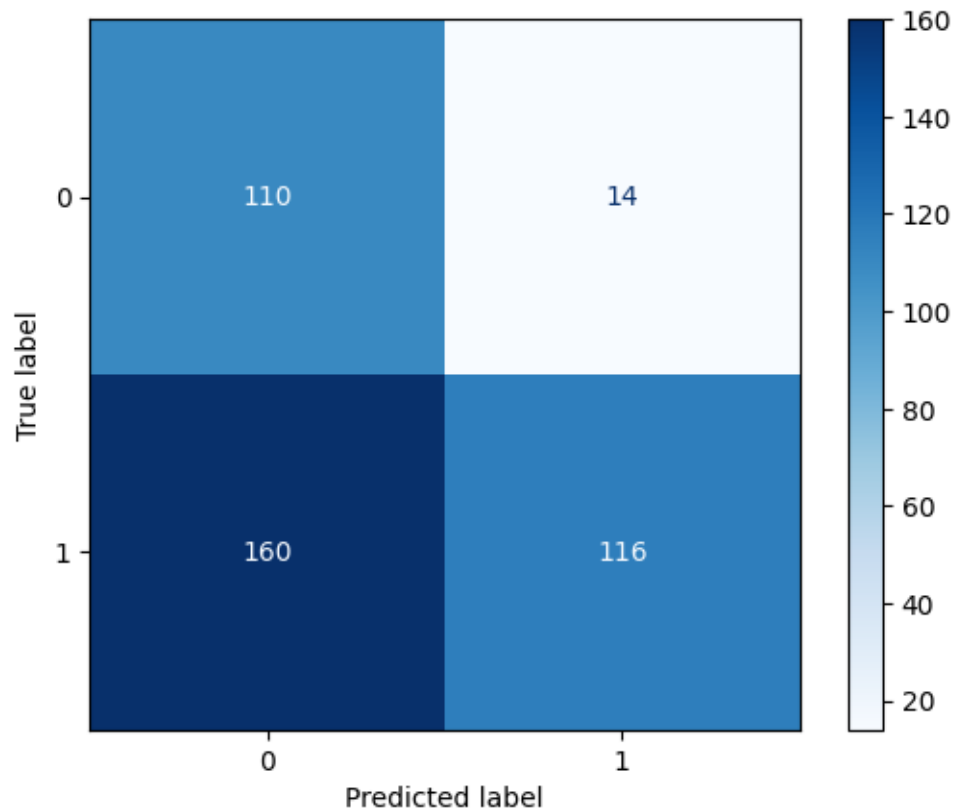
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show();
print('Confusion Matrix')
disp.plot(cmap='Blues', include_values=True)
plt.show();

print('-----')
print('Performance Measures')
print('-----')
print(f'Precision:, {precision}')
print(f'Recall:, {recall}')
print(f'F1 Score:, {f1}')
print(f'Accuracy Score: {scores}')
print('-----')
print('Hyperparameter Tuning Results')
print('-----')
print(f'Best parameters:, {grid_search.best_params_}')
print(f'Best score:", {grid_search.best_score_}')
print(f'False Positive rate: {conf_mtx[0][1] / (conf_mtx[0][1]+conf_mtx[0][0])}')

```



Confusion Matrix



Performance Measures

Precision:, 0.8923076923076924

Recall:, 0.42028985507246375

F1 Score:, 0.5714285714285715

Accuracy Score: 0.565

Hyperparameter Tuning Results

Best parameters:, {'var_smoothing': 0.0001}

Best score:", 0.5459604519774012

False Positive rate: 0.11290322580645161

Observation

- The precision on this model is very good but a bit “off-balance” with the Recall score, we are trying to minimize the False Positives and this model tends to do that as well. The problem can arise from the imbalanced data.

0.3.1 Simulating Best Score

```
[45]: i = 0

n_b_train_score = []
n_b_test_score = []
diff = []

while i < 200:

    # Splitting
    x_train, x_test, y_train, y_test = train_test_split(x, y, train_size=.60)

    # Standardizing
    scaler = StandardScaler()
    xtrain_scale = scaler.fit_transform(x_train)
    xtest = scaler.transform(x_test)

    # Hyperparameter Tuning
    param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
    grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
    ↪cv=10, scoring='accuracy')
    grid_search.fit(xtrain_scale, y_train)

    # Training
    n_b_best = grid_search.best_estimator_
    n_b_best.fit(xtrain_scale, y_train)
    predictions = n_b_best.predict(xtest)

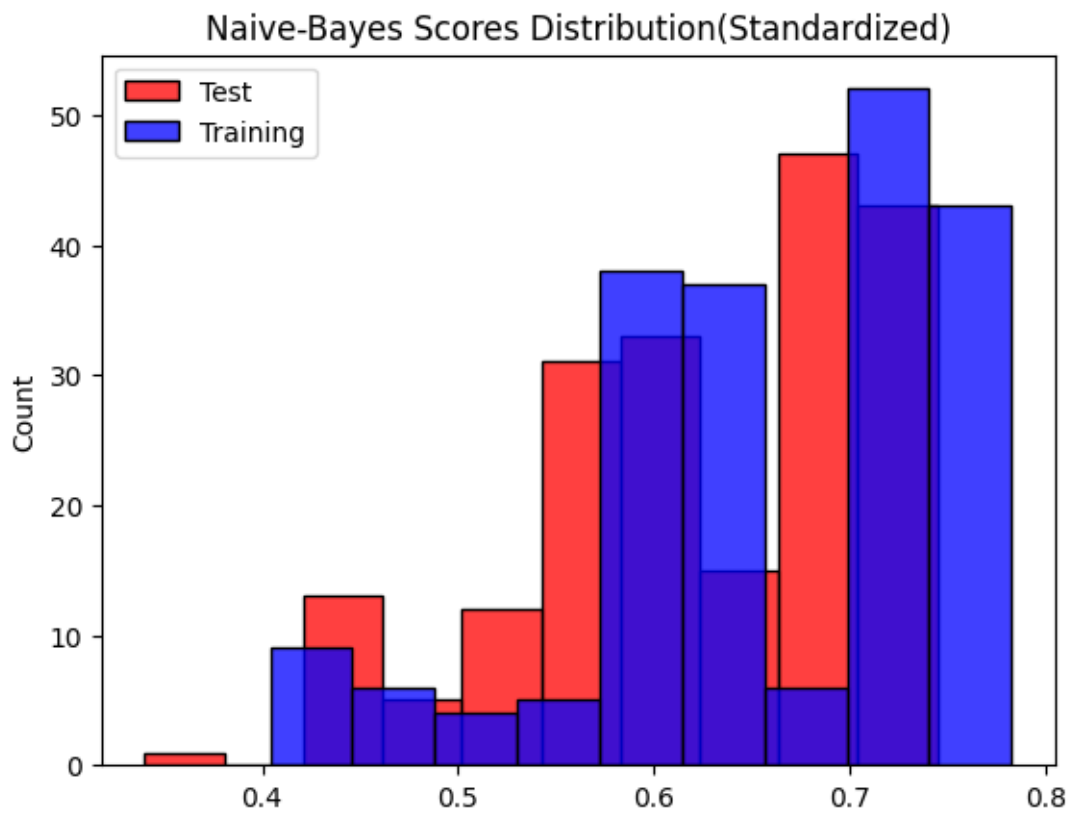
    # Performance
    train_score = n_b_best.score(xtrain_scale, y_train)
    test_score = n_b_best.score(xtest, y_test)
    n_b_train_score.append(train_score)
    n_b_test_score.append(test_score)
    diff_Score = train_score - test_score

    diff.append(diff_Score)

    i +=1

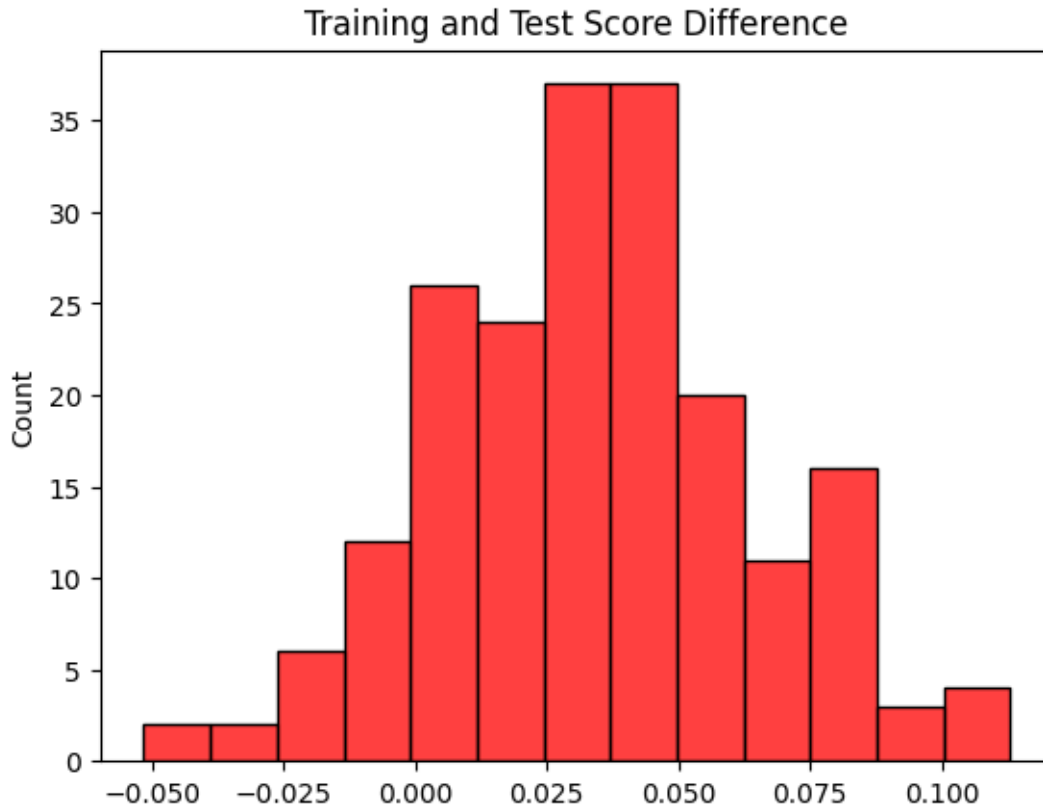
sns.histplot(data= n_b_test_score, color = 'red', label = 'Test')
sns.histplot(data= n_b_train_score, color = 'blue', label = 'Training')
plt.title('Naive-Bayes Scores Distribution(Standardized)')
plt.legend()
plt.show();
print(f'Mean Train Score: {np.mean(n_b_train_score)}')
```

```
print(f'Mean Test Score: {np.mean(n_b_test_score)}')  
print(f'Mean Diff score: {np.mean(diff)}')
```



Mean Train Score: 0.6587145242070116
Mean Test Score: 0.6248874999999999
Mean Diff score: 0.03382702420701169

```
[46]: sns.histplot(data= diff, color = 'red')  
plt.title('Training and Test Score Difference')  
plt.show();
```



Observation - From the simulations we can see that our score is very large spread and it is skewed to the left, the mean difference of our score is very close to zero indicating that our model is performing well, but like I mentioned before, it could be that the model is not learning well due to the imbalanced dataset, the spread is quite high as well and the difference of score having a negative means that the model is performing better on the test set rather than the training set which is preferable.

0.4 Naive Bayes Algorithm with Balanced Target and Standardized

```
[47]: #Splitting
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= .60,
↳ random_state= 42)

#Balancing
smote = SMOTE()
x_train_, y_train_ = smote.fit_resample(x_train,y_train)

#Standardizing
scaler = StandardScaler()
xtrain_scale = scaler.fit_transform(x_train_)
xtest = scaler.transform(x_test)
```



```

# Hyperparameter Tuning
param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
    ↪cv=10, scoring='accuracy')
grid_search.fit(xtrain_scale, y_train_)

# Training
n_b_best = grid_search.best_estimator_
n_b_best.fit(xtrain_scale, y_train_)
predictions= n_b_best.predict(xtest)

#Performance
scores = n_b_best.score(xtest, y_test)
conf_mtrx = confusion_matrix(y_test, predictions)
precision = precision_score(y_test, predictions)
recall = recall_score(y_test, predictions)
f1 = f1_score(y_test, predictions)
disp = ConfusionMatrixDisplay(confusion_matrix=conf_mtrx)

#Plotting ROC Curve
y_pred_proba = n_b_best.predict_proba(xtest)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba, pos_label=1)
roc_auc = auc(fpr, tpr)

#Results
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve(Standardized)
    ↪(area = {round(roc_auc,3)})')
plt.fill_between(fpr, tpr, color='gray', alpha=0.3)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show();
print('Confusion Matrix')
disp.plot(cmap='Blues', include_values=True)
plt.show();

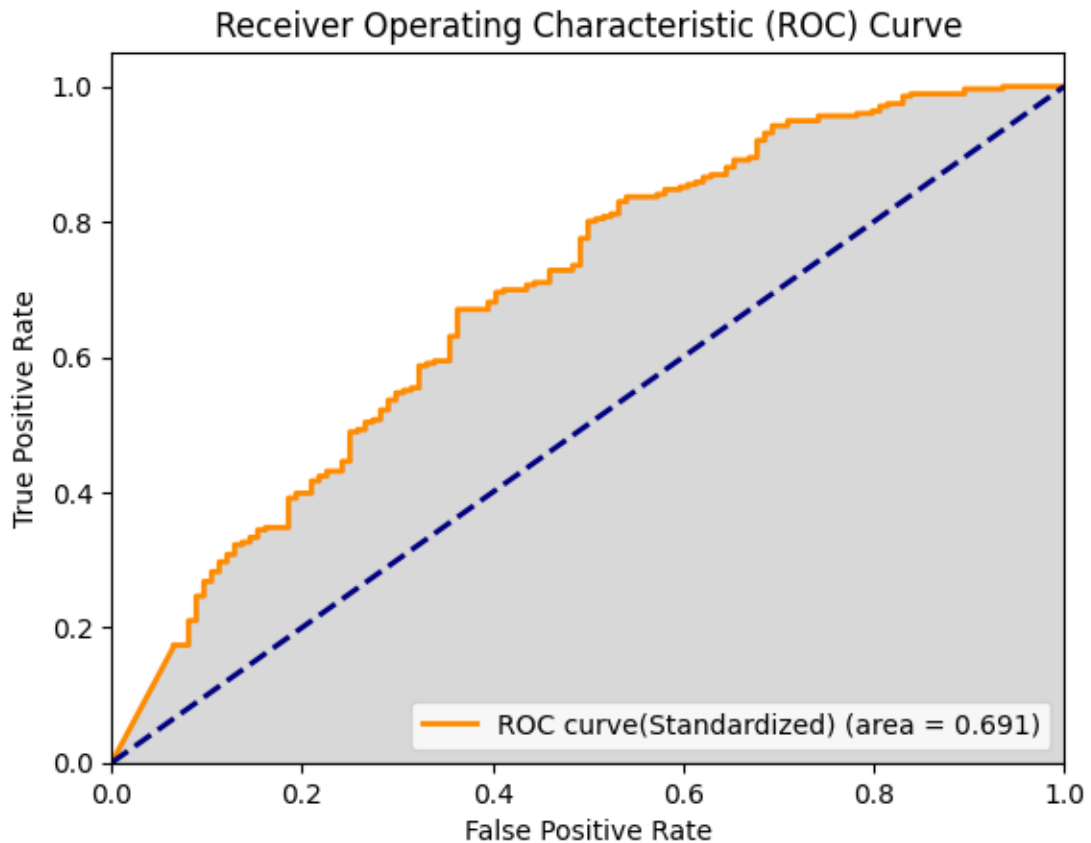
print('-----')
print('Performance Measures')
print('-----')
print(f'Precision:, {precision}')
print(f'Recall:, {recall}')

```

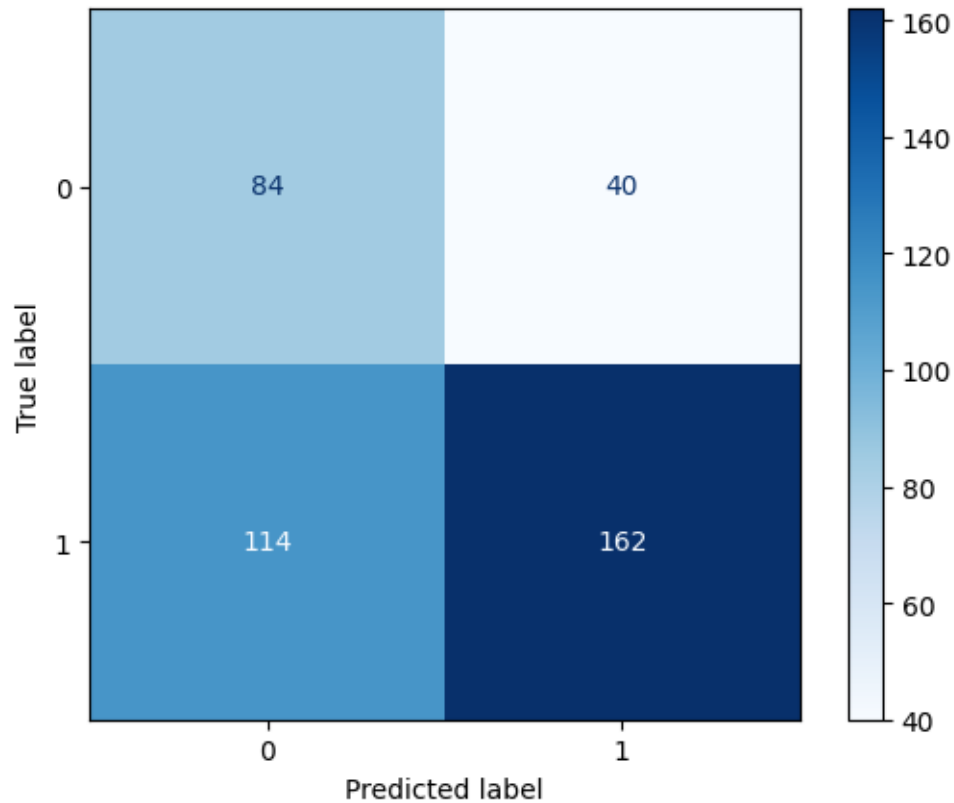
```

print(f'F1 Score:, {f1}')
print(f'Accuracy Score: {scores}')
print('-----')
print('Hyperparameter Tuning Results')
print('-----')
print(f'Best parameters:, {grid_search.best_params_}')
print(f'Best score:", {grid_search.best_score_}')
print(f'False Positive rate: {conf_mtx[0][1] / (
    ↪(conf_mtx[0][1]+conf_mtx[0][0])}')

```



Confusion Matrix



Performance Measures

Precision:, 0.801980198019802
Recall:, 0.5869565217391305
F1 Score:, 0.6778242677824268
Accuracy Score: 0.615

Hyperparameter Tuning Results

Best parameters:, {'var_smoothing': 0.0001}
Best score:", 0.7119327731092436
False Positive rate: 0.3225806451612903

Observations

- Precision has dropped but recall has increased, this means that our model has a good proportion of having large true positives while also making correct predictions. Our False Positive Rate has increased but it is balanced with the Recall score. The AUC is near the 50/50 threshold, it could potentially cause issues if our model is slightly “guessing” the predictions.

0.4.1 Simulating Best Score

```
[48]: i = 0
n_b_train_score = []
n_b_test_score = []
diff = []

while i < 200:

    #Splitting
    x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= .60)

    #Balancing
    smote = SMOTE()
    x_train_, y_train_ = smote.fit_resample(x_train,y_train)

    #Standardizing
    scaler = StandardScaler()
    xtrain_scale = scaler.fit_transform(x_train_)
    xtest = scaler.transform(x_test)

    # Hyperparameter Tuning
    param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
    grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
    ↪cv=10, scoring='accuracy')
    grid_search.fit(xtrain_scale, y_train_)

    # Training
    n_b_best = grid_search.best_estimator_
    n_b_best.fit(xtrain_scale, y_train_)
    predictions = n_b_best.predict(xtest)

    # Performance
    train_score = n_b_best.score(xtrain_scale, y_train_)
    test_score = n_b_best.score(xtest, y_test)
    n_b_train_score.append(train_score)
    n_b_test_score.append(test_score)
    diff_Score = train_score - test_score

    diff.append(diff_Score)

    i +=1

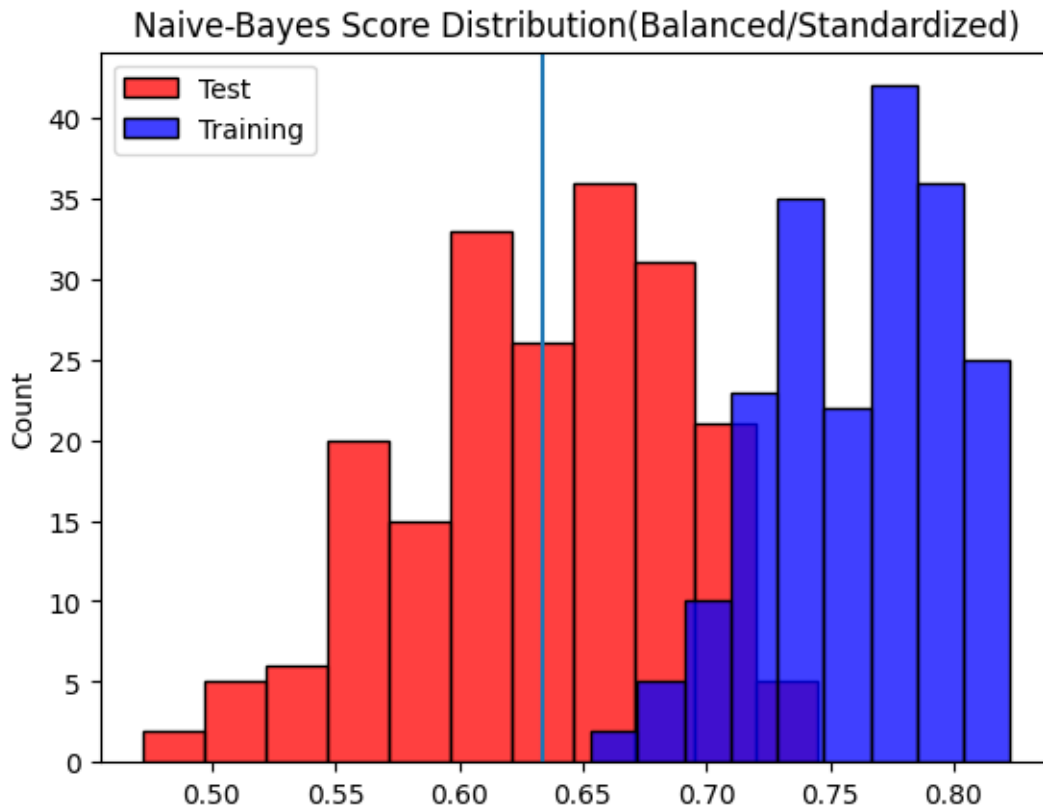
sns.histplot(data= n_b_test_score, color = 'red', label = 'Test')
```

```

sns.histplot(data= n_b_train_score, color = 'blue', label = 'Training')
plt.title('Naive-Bayes Score Distribution(Balanced/Standardized)')
plt.legend()
plt.axvline(np.mean(n_b_test_score))
plt.show();

print(f'Mean Train Score: {np.mean(n_b_train_score)}')
print(f'Mean Test Score: {np.mean(n_b_test_score)}')
print(f'Mean Diff score: {np.mean(diff)}')

```

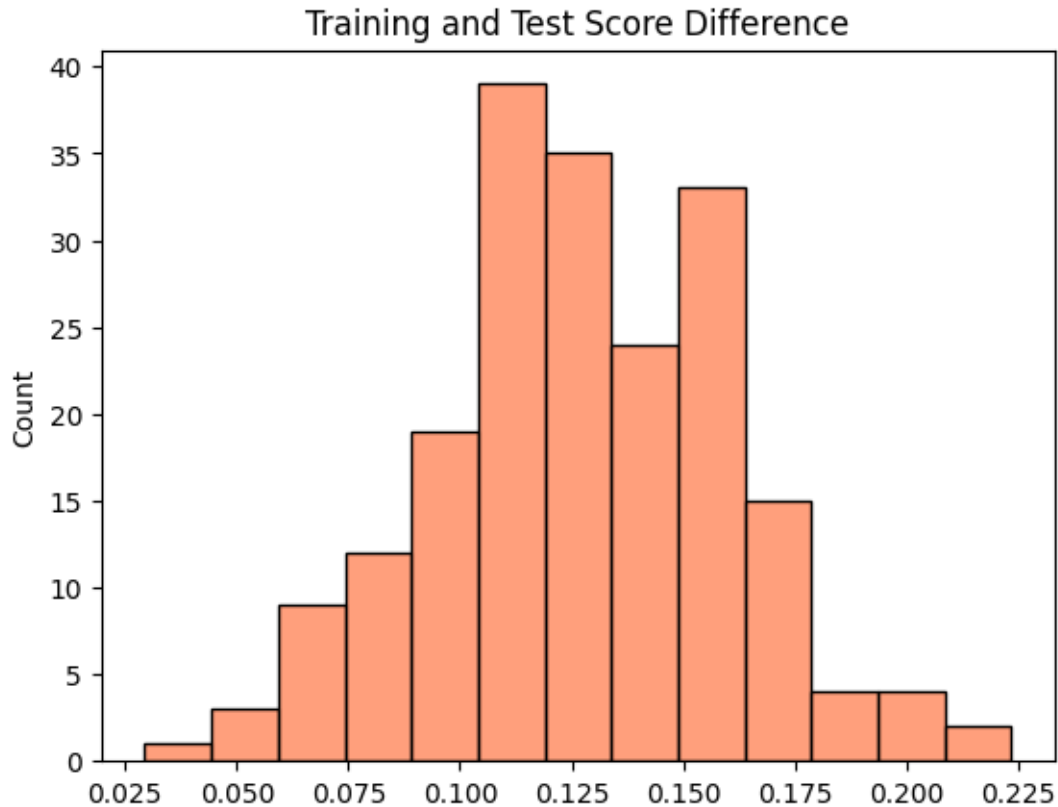


Mean Train Score: 0.7613419087298713
 Mean Test Score: 0.633625
 Mean Diff score: 0.12771690872987132

```

[49]: sns.histplot(data = diff, color = 'coral')
plt.title('Training and Test Score Difference')
plt.show();

```



Observation - In comparison to the previous simulation, this model is doing better on the training data, which could also be a sign of overfitting.

0.5 Without modifications to the dataset

```
[50]: #Splitting
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= .60,
    random_state= 42)

# Hyperparameter Tuning
param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
    cv=10, scoring='accuracy')
grid_search.fit(x_train, y_train)

# Training
n_b_best = grid_search.best_estimator_
n_b_best.fit(x_train, y_train)
predictions_ = n_b_best.predict(x_test)

#Performance
```

```

scores = n_b_best.score(x_test, y_test)
conf_mtx = confusion_matrix(y_test, predictions_)
precision = precision_score(y_test, predictions_)
recall = recall_score(y_test, predictions_)
f1 = f1_score(y_test, predictions_)
disp = ConfusionMatrixDisplay(confusion_matrix=conf_mtx)

#Results
print('Confusion Matrix')
disp.plot(cmap='Blues', include_values=True)
plt.show();

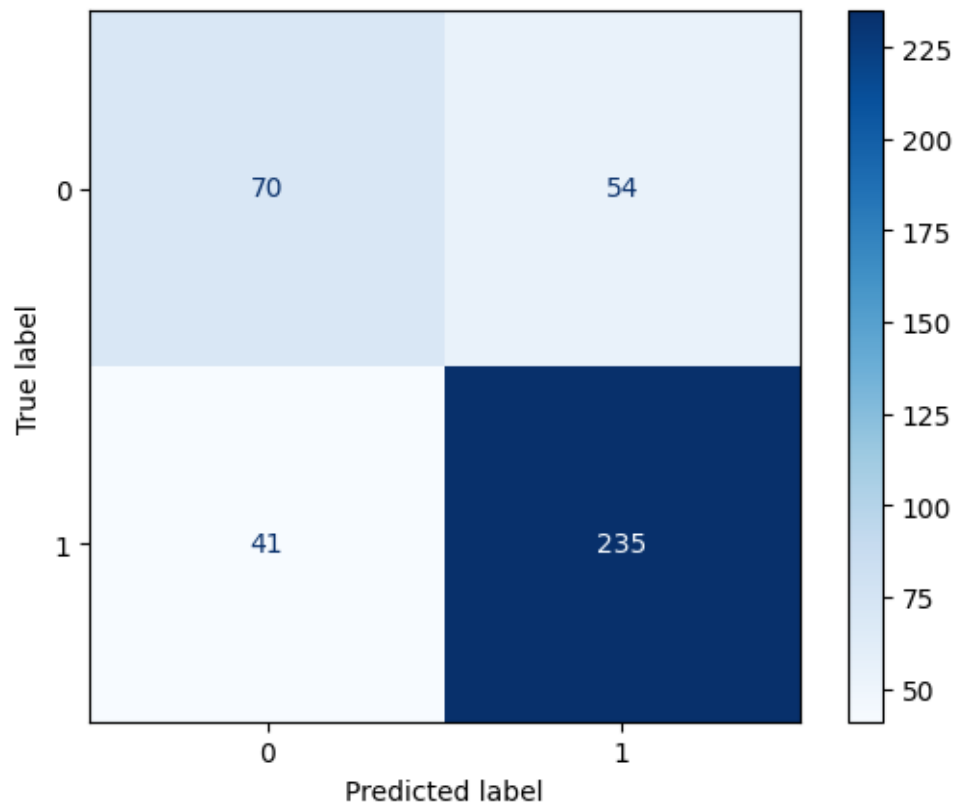
#Plotting ROC Curve
y_pred_proba = n_b_best.predict_proba(x_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba, pos_label=1)
roc_auc = auc(fpr, tpr)

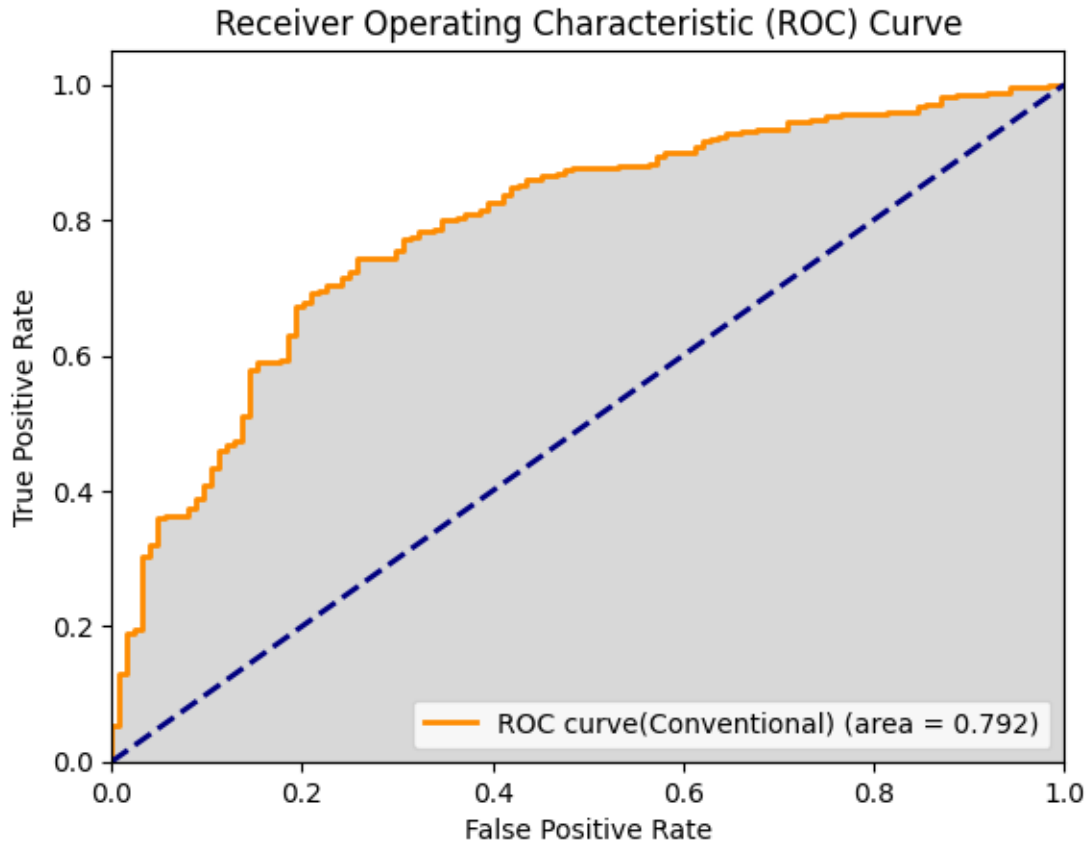
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve(Conventional)')
↳(area = {round(roc_auc,3)})')
plt.fill_between(fpr, tpr, color='gray', alpha=0.3)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show();

print('-----')
print('Performance Measures')
print('-----')
print(f'Precision:, {precision}')
print(f'Recall:, {recall}')
print(f'F1 Score:, {f1}')
print(f'Accuracy Score: {scores}')
print('-----')
print('Hyperparameter Tuning Results')
print('-----')
print(f'Best parameters:, {grid_search.best_params_}')
print(f'Best score:", {grid_search.best_score_}')
print(f'False Positive rate: {conf_mtx[0][1] / }
↳(conf_mtx[0][1]+conf_mtx[0][0]))')

```

Confusion Matrix





Performance Measures

Precision:, 0.8131487889273357
 Recall:, 0.8514492753623188
 F1 Score:, 0.831858407079646
 Accuracy Score: 0.7625

Hyperparameter Tuning Results

Best parameters:, {'var_smoothing': 1e-08}
 Best score:", 0.7495762711864407
 False Positive rate: 0.43548387096774194

Observations

- We got higher results by leaving the model intact, no standardization and no balance dataset. Our AUC is close to .80, a percentage that can be categorized as a good performing model. The ratio of precision and recall or F1 score is above .80 which means that our model is good a predicting a high number of True Positives without sacrificing its accuracy. Overall this model can be a candidate.

0.5.1 Simulating Best Score

```
[51]: i = 0
n_b_train_score = []
n_b_test_score = []
diff = []

while i < 200:

    #Splitting
    x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= .60)

    # Hyperparameter Tuning
    param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
    grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
    ↪cv=10, scoring='accuracy')
    grid_search.fit(x_train, y_train)

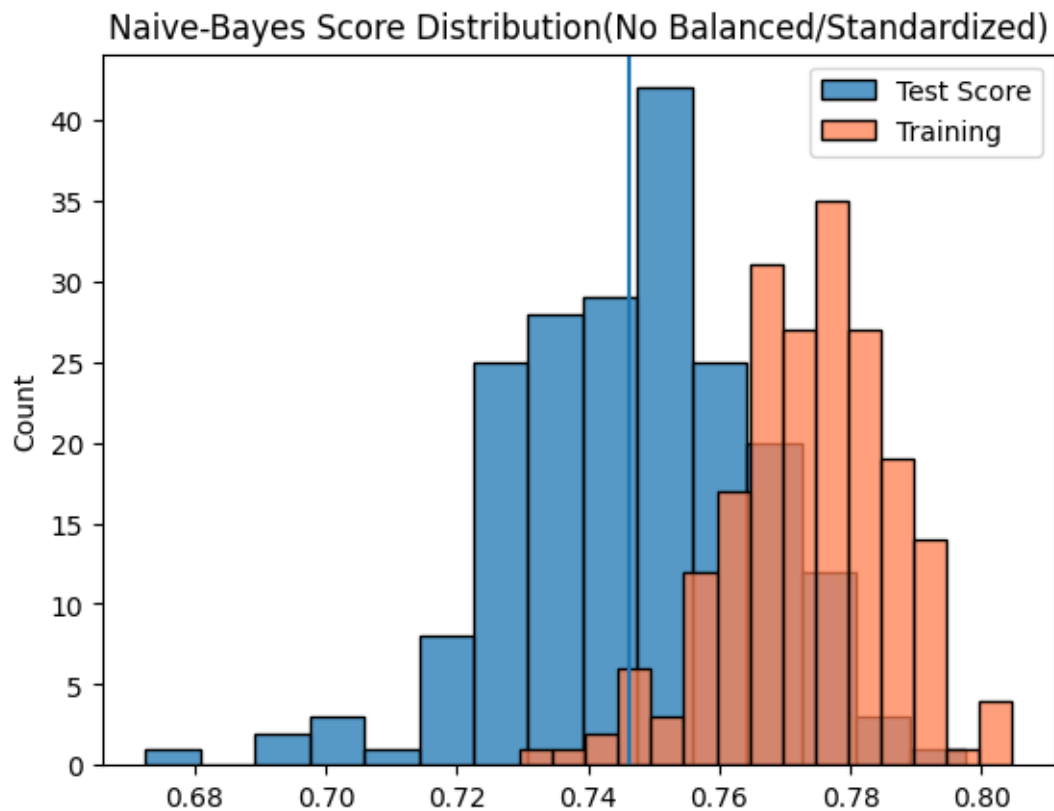
    # Training
    n_b_best = grid_search.best_estimator_
    n_b_best.fit(x_train, y_train)
    predictions = n_b_best.predict(x_test)

    # Performance
    train_score = n_b_best.score(x_train, y_train)
    test_score = n_b_best.score(x_test, y_test)
    n_b_train_score.append(train_score)
    n_b_test_score.append(test_score)
    diff_score = train_score - test_score
    diff.append(diff_score)

    i +=1

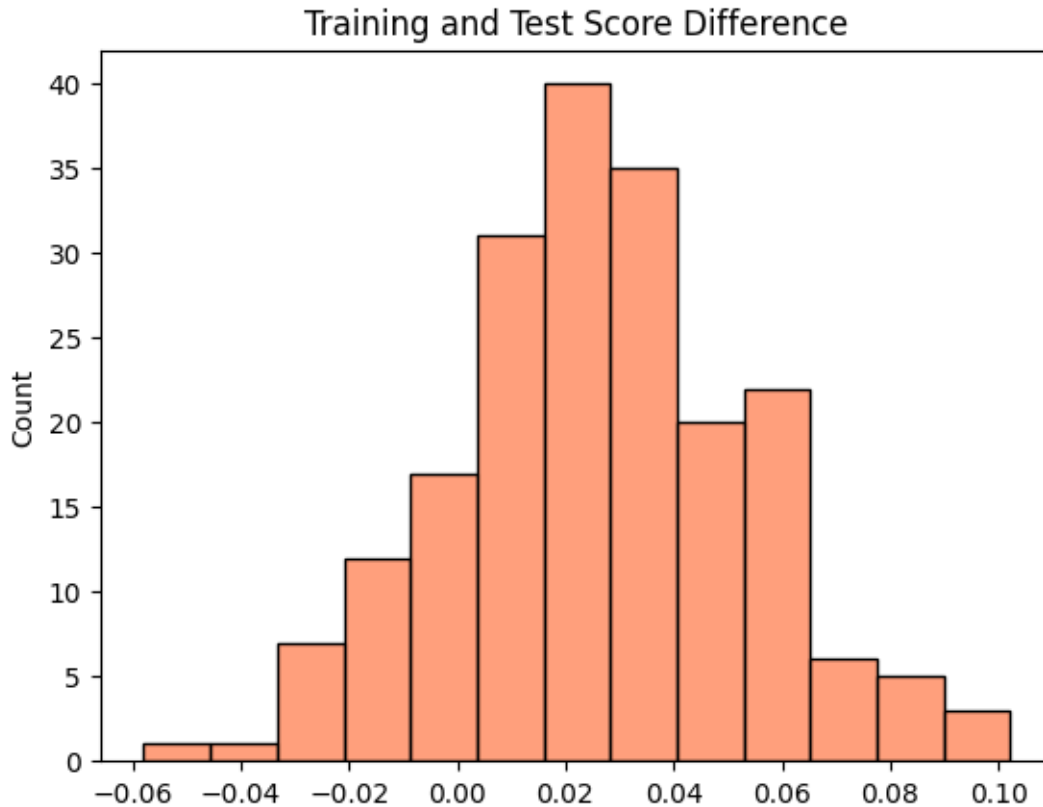
sns.histplot(data=n_b_test_score, bins=15, label = 'Test Score')
sns.histplot(data = n_b_train_score, bins = 15, color = 'coral', label =
    ↪'Training')
plt.title('Naive-Bayes Score Distribution(No Balanced/Standardized)')
plt.legend()
plt.axvline(np.mean(n_b_test_score))
plt.show();

print(f'Mean Train Score: {np.mean(n_b_train_score)}')
print(f'Mean Test Score: {np.mean(n_b_test_score)}')
print(f'Mean Diff score: {np.mean(diff)}')
```



Mean Train Score: 0.7726794657762938
Mean Test Score: 0.7461500000000001
Mean Diff score: 0.026529465776293823

```
[52]: sns.histplot(data = diff, color = 'coral')  
plt.title('Training and Test Score Difference')  
plt.show();
```



The difference in scores has a range with negative numbers, again, it is an observation that our models is performing better on our test set rather than our training set.

0.6 With Balanced Target And No Standardization

```
[53]: #Splitting
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= .60,
↳random_state=42)

#Balancing
smote = SMOTE()
x_train_, y_train_ = smote.fit_resample(x_train,y_train)

# Hyperparameter Tuning
param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
↳cv=10, scoring='accuracy')
grid_search.fit(x_train_, y_train_)

# Training
```

```

n_b_best = grid_search.best_estimator_
n_b_best.fit(x_train_, y_train_)
predictions_ = n_b_best.predict(x_test)

#Performance
scores = n_b_best.score(x_test, y_test)
conf_mtrx = confusion_matrix(y_test, predictions_)
precision = precision_score(y_test, predictions_)
recall = recall_score(y_test, predictions_)
f1 = f1_score(y_test, predictions_)
disp = ConfusionMatrixDisplay(confusion_matrix=conf_mtrx)

#Results
print('Confusion Matrix')
disp.plot(cmap='Blues', include_values=True)
plt.show();

#Plotting ROC Curve
y_pred_proba = n_b_best.predict_proba(x_test)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba, pos_label=1)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve(Balanced) (area_
    ↳ = {round(roc_auc,3)})')
plt.fill_between(fpr, tpr, color='gray', alpha=0.3)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show();

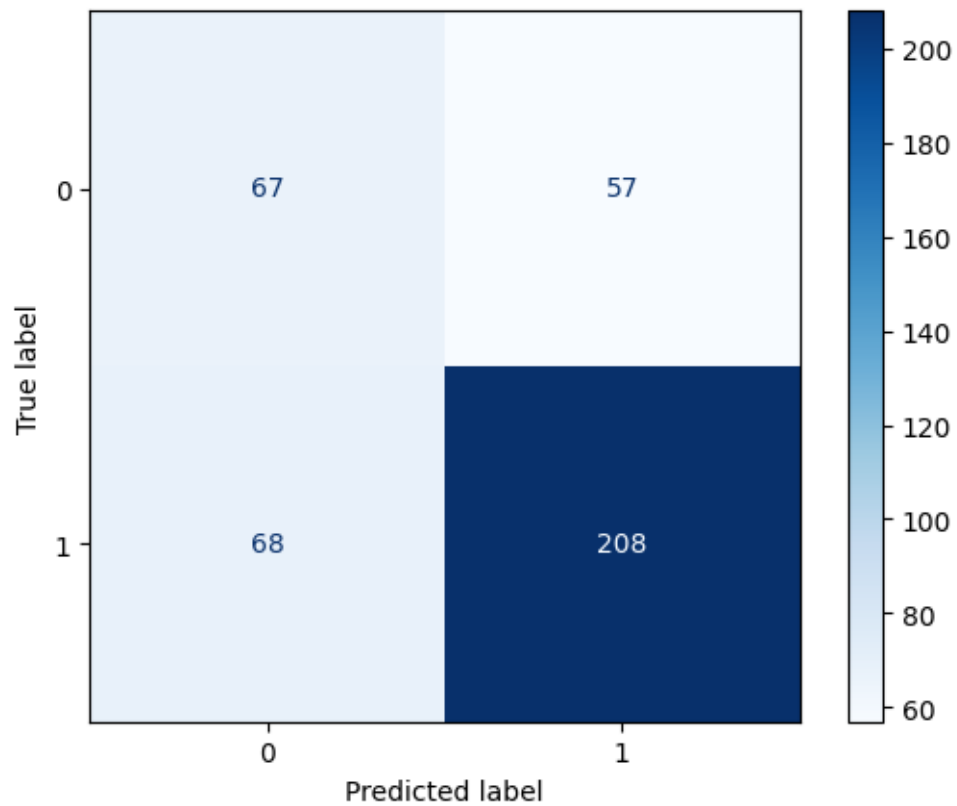
print('-----')
print('Performance Measures')
print('-----')
print(f'Precision:, {precision}')
print(f'Recall:, {recall}')
print(f'F1 Score:, {f1}')
print(f'Accuracy Score: {scores}')
print('-----')
print('Hyperparameter Tuning Results')
print('-----')

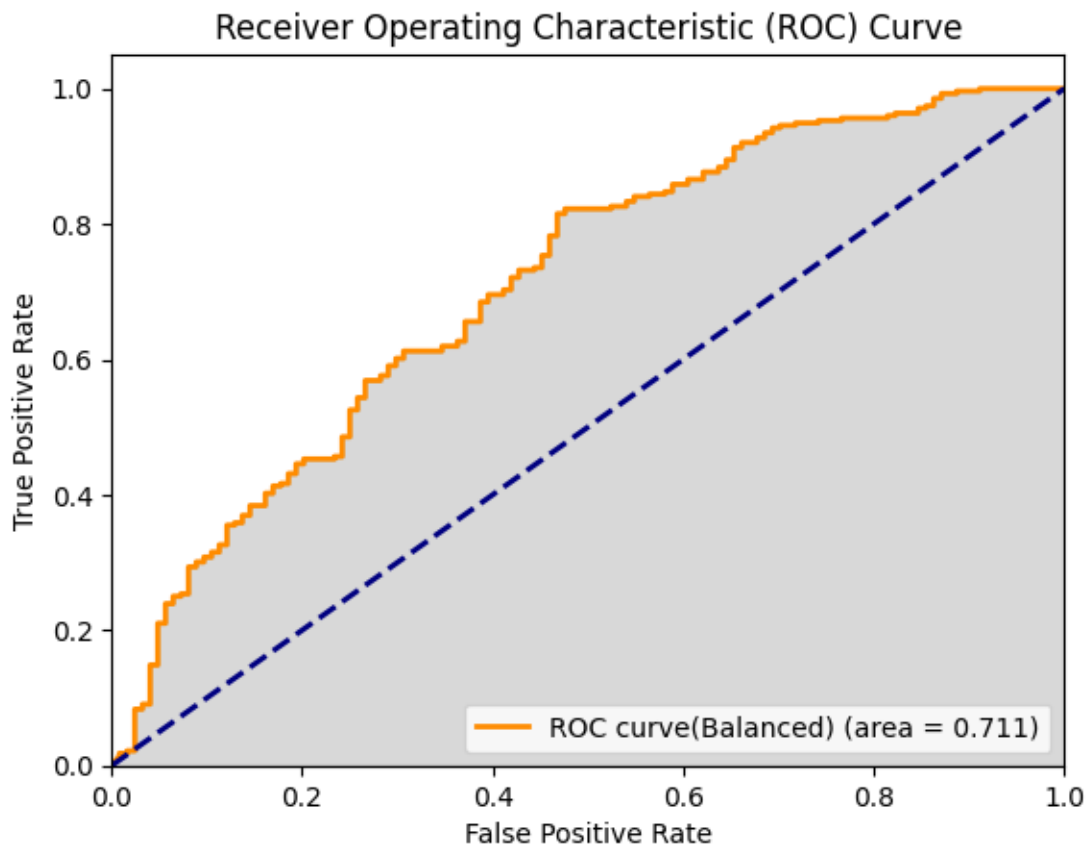
```

```
print(f'Best parameters:', {grid_search.best_params_})  
print(f'Best score:', {grid_search.best_score_})  
print(f'False Positive rate: {conf_mtx[0][1] /  
↪ (conf_mtx[0][1]+conf_mtx[0][0])}')  

```

Confusion Matrix





Performance Measures

Precision:, 0.7849056603773585
 Recall:, 0.7536231884057971
 F1 Score:, 0.7689463955637708
 Accuracy Score: 0.6875

Hyperparameter Tuning Results

Best parameters:, {'var_smoothing': 1e-09}
 Best score:", 0.7770448179271708
 False Positive rate: 0.4596774193548387

This model performed very similar to the conventional algorithm except the score were lower but that ratio is consistent, give then results from the previous model we can discard this one.

0.6.1 Simulating Best Score

```
[54]: i = 0
n_b_score = []
n_b_train_score = []
n_b_test_score = []
diff = []

while i < 200:

    #Splitting
    x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= .60)

    #Balancing
    smote = SMOTE()
    x_train_, y_train_ = smote.fit_resample(x_train,y_train)

    # Hyperparameter Tuning
    param_grid = {'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4]}
    grid_search = GridSearchCV(estimator=GaussianNB(), param_grid=param_grid,
    ↪cv=10, scoring='accuracy')
    grid_search.fit(x_train_, y_train_)

    # Training
    n_b_best = grid_search.best_estimator_
    n_b_best.fit(x_train_, y_train_)
    predictions_ = n_b_best.predict(x_test)

    # Performance
    train_score = n_b_best.score(x_train, y_train)
    test_score = n_b_best.score(x_test, y_test)
    n_b_train_score.append(train_score)
    n_b_test_score.append(test_score)
    diff_score = train_score - test_score
    diff.append(diff_score)

    i +=1

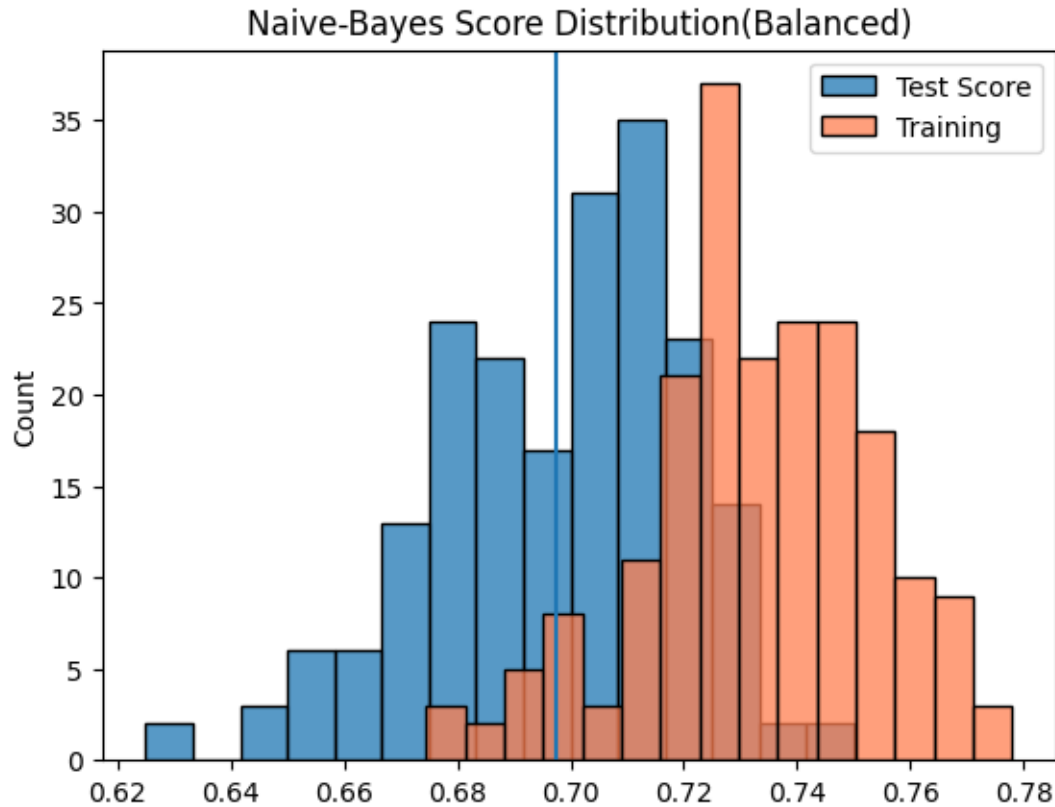
sns.histplot(data=n_b_test_score, bins=15, label = 'Test Score')
sns.histplot(data = n_b_train_score, bins = 15, color = 'coral', label =
    ↪'Training')
plt.title('Naive-Bayes Score Distribution(Balanced)')
plt.legend()
plt.axvline(np.mean(n_b_test_score))
plt.show();
```



```

print(f'Mean Train Score: {np.mean(n_b_train_score)}')
print(f'Mean Test Score: {np.mean(n_b_test_score)}')
print(f'Mean Diff score: {np.mean(diff)}')

```

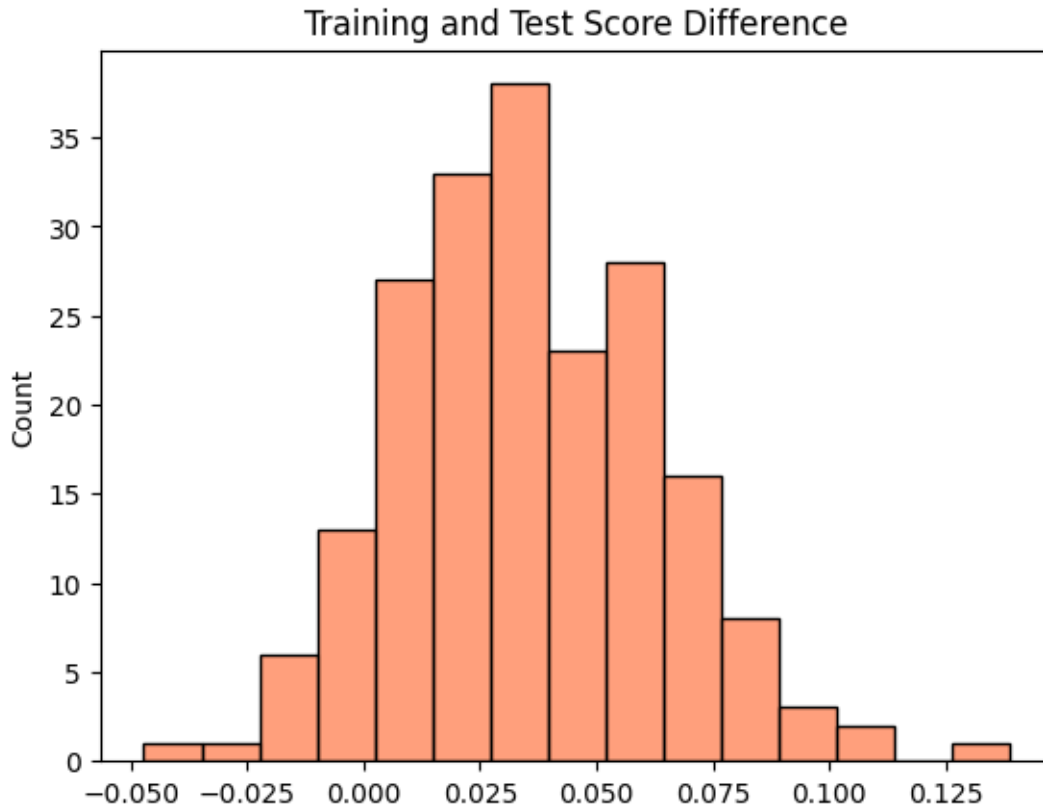


Mean Train Score: 0.7331218697829717
 Mean Test Score: 0.6975625
 Mean Diff score: 0.035559369782971625

```

[55]: sns.histplot(data = diff, color = 'coral')
plt.title('Training and Test Score Difference')
plt.show();

```



The abrupt cut in the histogram to negative values means that the model can have a high level of variability in its results. Although in some instances it performs better in the test set, the count tends to be low therefore the model can have problems adapting to new data.

0.7 Optimized Gradient Boosting Machine Implementation for Balanced Classification

```
[56]: # Splitting
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size=0.75)

# Balancing
smote = SMOTE()
X_resampled, Y_resampled = smote.fit_resample(x_train, y_train)

# Hyperparameter Tuning for Gradient Boosting
param_grid = {
    'n_estimators': [100, 200, 300], # Number of boosting stages to be run
    'learning_rate': [0.01, 0.1, 0.2], # Shrinks the contribution of each tree
    'by_learning_rate': 0.01
}
```

```

    'max_depth': [3, 4, 5] # Maximum depth of the individual regression
    ↪estimators
}
grid_search = GridSearchCV(estimator=GradientBoostingClassifier(),
    ↪param_grid=param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_resampled, Y_resampled)

# Training with the best parameters found
gbm_best = grid_search.best_estimator_
gbm_best.fit(X_resampled, Y_resampled)
predictions = gbm_best.predict(x_test)

# Performance evaluation
scores = gbm_best.score(x_test, y_test)
conf_mtrx = confusion_matrix(y_test, predictions)
precision = precision_score(y_test, predictions)
recall = recall_score(y_test, predictions)
f1 = f1_score(y_test, predictions)

# Confusion Matrix Display
disp = ConfusionMatrixDisplay(confusion_matrix=conf_mtrx)
print('Confusion Matrix')
disp.plot(cmap='Blues', include_values=True)
plt.show()

# Plotting ROC Curve
y_pred_proba = gbm_best.predict_proba(x_test)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba, pos_label=1)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (GBM) (area =
    ↪{round(roc_auc, 3)})')
plt.fill_between(fpr, tpr, color='gray', alpha=0.3)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

print('-----')
print('Performance Measures')
print('-----')
print(f'Precision: {precision}')

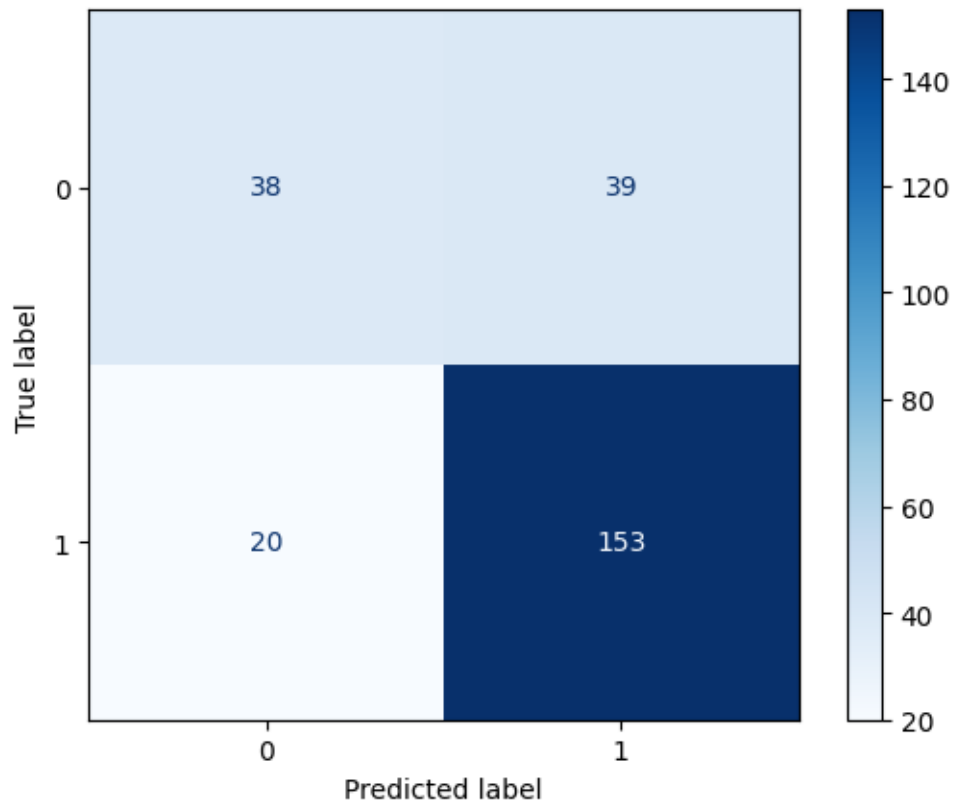
```

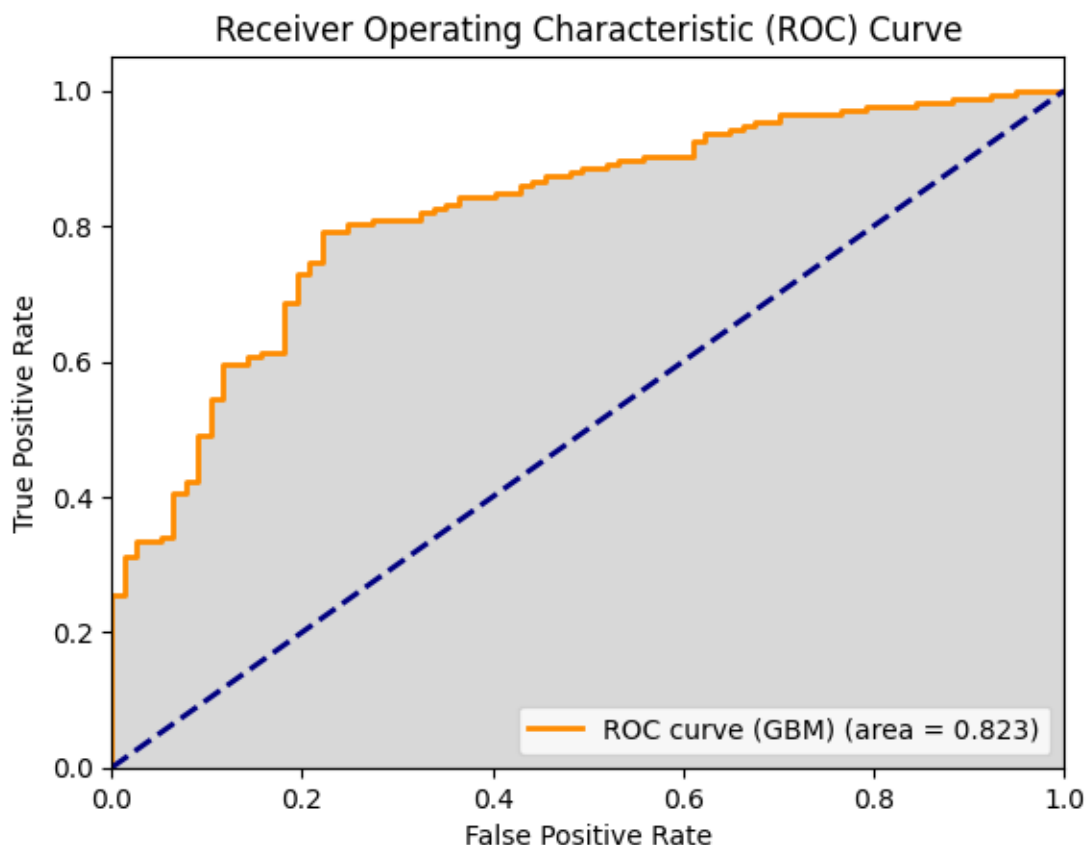
```

print(f'Recall: {recall}')
print(f'F1 Score: {f1}')
print(f'Accuracy Score: {scores}')
print('-----')
print('Hyperparameter Tuning Results')
print('-----')
print(f'Best parameters: {grid_search.best_params_}')
print(f'Best score: {grid_search.best_score_}')
print(f'False Positive rate: {conf_mtx[0][1] / (conf_mtx[0][1]+conf_mtx[0][0])}')

```

Confusion Matrix





Performance Measures

Precision: 0.796875
Recall: 0.884393063583815
F1 Score: 0.8383561643835618
Accuracy Score: 0.764

Hyperparameter Tuning Results

Best parameters: {'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 200}
Best score: 0.8025366734371474
False Positive rate: 0.5064935064935064

0.7.1 GBM Result Interpretation

The Model has a high rate of predicting true positives (Precision) and catching the majority of the total actual positives (Recall). The high F1 score indicates that the model effectively balances both metrics. The models ability to predict all True Positives and True Negatives is its weakest score, though it is still relatively high (Accuracy). Overall, the performance measures are well rounded.

```

[57]: # Initialize lists to store scores and differences
gbm_precision = []
gbm_recall = []
gbm_f1_score = []
gbm_accuracy = []
gbm_roc_auc = []
gbm_diff = []

# Counter
i = 0
# Number of iterations
iterations = 50

while i < iterations:
    # Data splitting
    x_train, x_test, y_train, y_test = train_test_split(x, y, train_size=0.60)

    # Data balancing with SMOTE
    smote = SMOTE()
    X_resampled, Y_resampled = smote.fit_resample(x_train, y_train)

    # Hyperparameter tuning for Gradient Boosting
    param_grid = {
        'n_estimators': [100, 200, 300],
        'learning_rate': [0.01, 0.1, 0.2],
        'max_depth': [3, 4, 5]
    }
    grid_search = GridSearchCV(estimator=GradientBoostingClassifier(),
    ↪ param_grid=param_grid, cv=10, scoring='accuracy')
    grid_search.fit(X_resampled, Y_resampled)

    # Training with the best parameters found
    gbm_best = grid_search.best_estimator_
    gbm_best.fit(X_resampled, Y_resampled)
    predictions = gbm_best.predict(x_test)

    # Performance evaluation
    precision = precision_score(y_test, predictions)
    recall = recall_score(y_test, predictions)
    f1_score_val = f1_score(y_test, predictions)
    accuracy = gbm_best.score(x_test, y_test)
    train_score = gbm_best.score(X_resampled, Y_resampled)

    # Difference in training and testing accuracy
    diff_score = train_score - accuracy

    # Store scores

```

```

gbm_precision.append(precision)
gbm_recall.append(recall)
gbm_f1_score.append(f1_score_val)
gbm_accuracy.append(accuracy)
gbm_diff.append(diff_score)

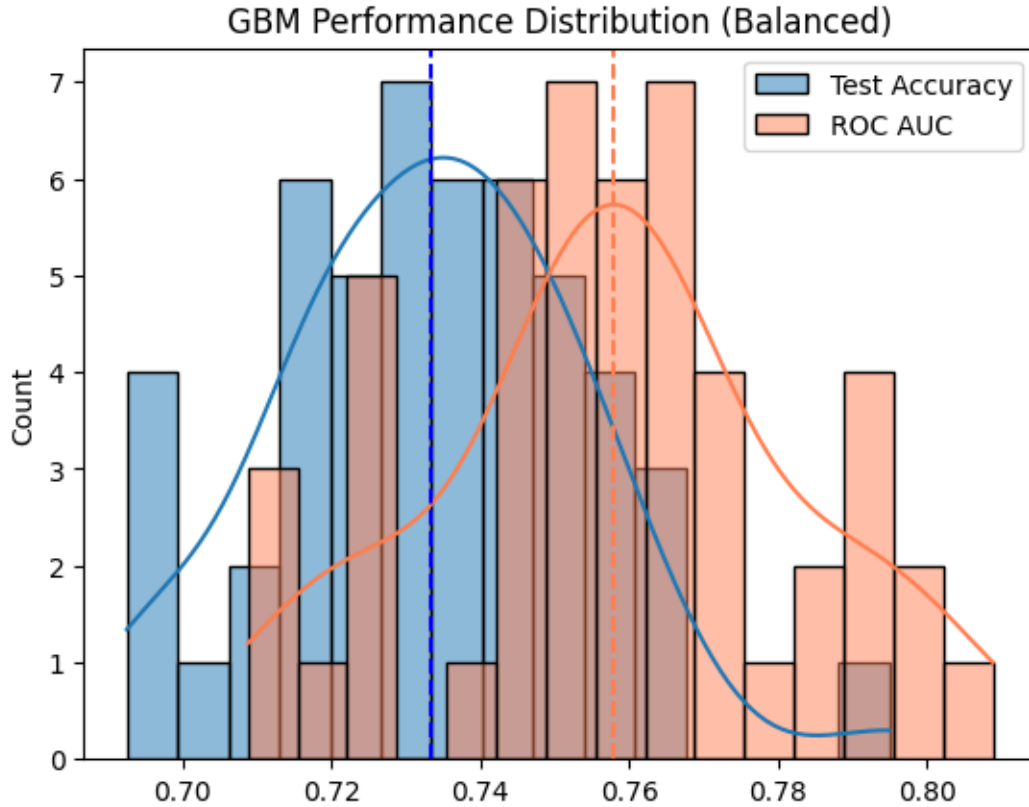
# ROC Curve calculations
y_pred_proba = gbm_best.predict_proba(x_test)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba, pos_label=1)
roc_auc = auc(fpr, tpr)
gbm_roc_auc.append(roc_auc)

i += 1

# Plotting the results
sns.histplot(data=gbm_accuracy, bins=15, label='Test Accuracy', kde=True)
sns.histplot(data=gbm_roc_auc, bins=15, color='coral', label='ROC AUC',
             ↪kde=True)
plt.title('GBM Performance Distribution (Balanced)')
plt.legend()
plt.axvline(x=np.mean(gbm_accuracy), color='blue', linestyle='--',
            ↪label='Average Test Accuracy')
plt.axvline(x=np.mean(gbm_roc_auc), color='coral', linestyle='--',
            ↪label='Average ROC AUC')
plt.show()

# Print the average scores
print(f'Mean Precision: {np.mean(gbm_precision)}')
print(f'Mean Recall: {np.mean(gbm_recall)}')
print(f'Mean F1 Score: {np.mean(gbm_f1_score)}')
print(f'Mean Accuracy: {np.mean(gbm_accuracy)}')
print(f'Mean ROC AUC: {np.mean(gbm_roc_auc)}')
print(f'Mean Difference Score: {np.mean(gbm_diff)}')

```



Mean Precision: 0.7979524838236869
Mean Recall: 0.830131521575474
Mean F1 Score: 0.8132233322883425
Mean Accuracy: 0.7333500000000001
Mean ROC AUC: 0.7579680046709298
Mean Difference Score: 0.24569384029906613

0.7.2 GBM Iteration Interpretation

The mean of the precision, recall, F1 Score, and Accuracy were just a tad below the single iteration results, with the weakest measurement (Accuracy) having the greatest difference. The mean ROC AUC score from the iterations was 0.758, substantiating the model's discriminatory ability. The small mean difference score of 0.246 between the training and testing accuracy further reinforces confidence in the model's reliability. These results indicate that GBM is capable of delivering reliable and consistent predictions.

0.8 RandomForest Algorithm (Resetted df)

We decided for the final algorithm, to reset the df and compare the results of a comparable algorithm but with much less data cleaning and feature manipulation to the previous two algorithm results to see how much of an impact our previous cleaning had yielded.


```
[58]: # Loading dataset and setting column names

df = pd.read_csv("german.data", sep=r'\s+')
column_names = ['checking_account', 'duration_month', 'credit_history',
↳ 'credit_purpose', 'credit_amount', 'savings_account', 'present_employment', 'disposable_income_
↳
↳ 'status_sex', 'debtors', 'residence_since', 'property', 'age', 'other_installments',
↳ 'housing', 'credits_at_current_bank', 'job', 'dependants', 'telephone',
↳ 'foreign_worker', 'class']

df.columns = column_names
df.columns
```

```
[58]: Index(['checking_account', 'duration_month', 'credit_history',
'credit_purpose', 'credit_amount', 'savings_account',
'present_employment', 'disposable_income_percent', 'status_sex',
'debtors', 'residence_since', 'property', 'age', 'other_installments',
'housing', 'credits_at_current_bank', 'job', 'dependants', 'telephone',
'foreign_worker', 'class'],
dtype='object')
```

```
[59]: # Separating "Explanatory", and "Response Variables" into separate DataFrames.
# Then encoding categorical values and changed "Class" variable into a string
↳ since it does not represent a numerical value but is rather a category.
explanatory = df.drop(columns=['class'])
explanatory_dummies = pd.get_dummies(explanatory)
response = df['class'].astype(str).replace({'1': 'Good', '2': 'Bad'})
df.groupby(['class']).count()
```

```
[59]:
```

	checking_account	duration_month	credit_history	credit_purpose	\
class					
1	699	699	699	699	
2	300	300	300	300	

	credit_amount	savings_account	present_employment	\
class				
1	699	699	699	
2	300	300	300	

	disposable_income_percent	status_sex	debtors	residence_since	\
class					
1	699	699	699	699	
2	300	300	300	300	

	property	age	other_installments	housing	credits_at_current_bank	\
class						
1	699	699	699	699	699	

2	300	300	300	300	300
	job	dependants	telephone	foreign_worker	
class					
1	699	699	699	699	
2	300	300	300	300	

```
[60]: # Creating training and testing sets
x_train, x_test, y_train, y_test = train_test_split(explanatory_dummies,
    ↳response, test_size = 0.2)
y_train_flat = np.ravel(y_train, order='C')
y_test_flat = np.ravel(y_test, order='C')

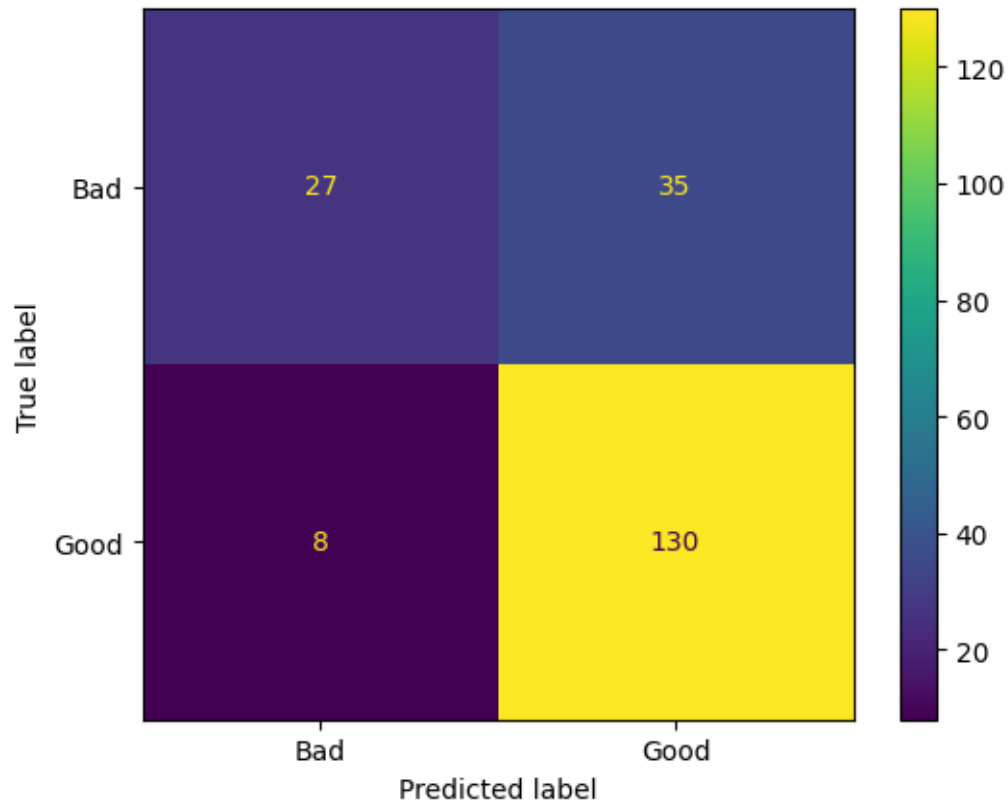
rf = RandomForestClassifier()
rf.fit(x_train, y_train_flat)

# Precision: Percent of "True Positives" identified correctly divided by all
    ↳predicted positives in dataset
# Recall: Percent of positive predictions were correct.
# I'd argue the data set is unbalanced
preds = rf.predict(x_test)
cr = classification_report(y_test, preds)
print(cr)

ConfusionMatrixDisplay.from_predictions(y_test, preds)
```

	precision	recall	f1-score	support
Bad	0.77	0.44	0.56	62
Good	0.79	0.94	0.86	138
accuracy			0.79	200
macro avg	0.78	0.69	0.71	200
weighted avg	0.78	0.79	0.76	200

```
[60]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x31896d6d0>
```



```
[61]: # Determining most important explanatory features
# Training Random Forest on new training data

feature_select = SelectFromModel(RandomForestClassifier(n_estimators = 50))
feature_select.fit(x_train, y_train)
features = feature_select.get_feature_names_out().tolist()

selected_features_train = x_train[features]
selected_features_test = x_test[features]

rf.fit(selected_features_train, y_train_flat)

# New classification report based on selected features.

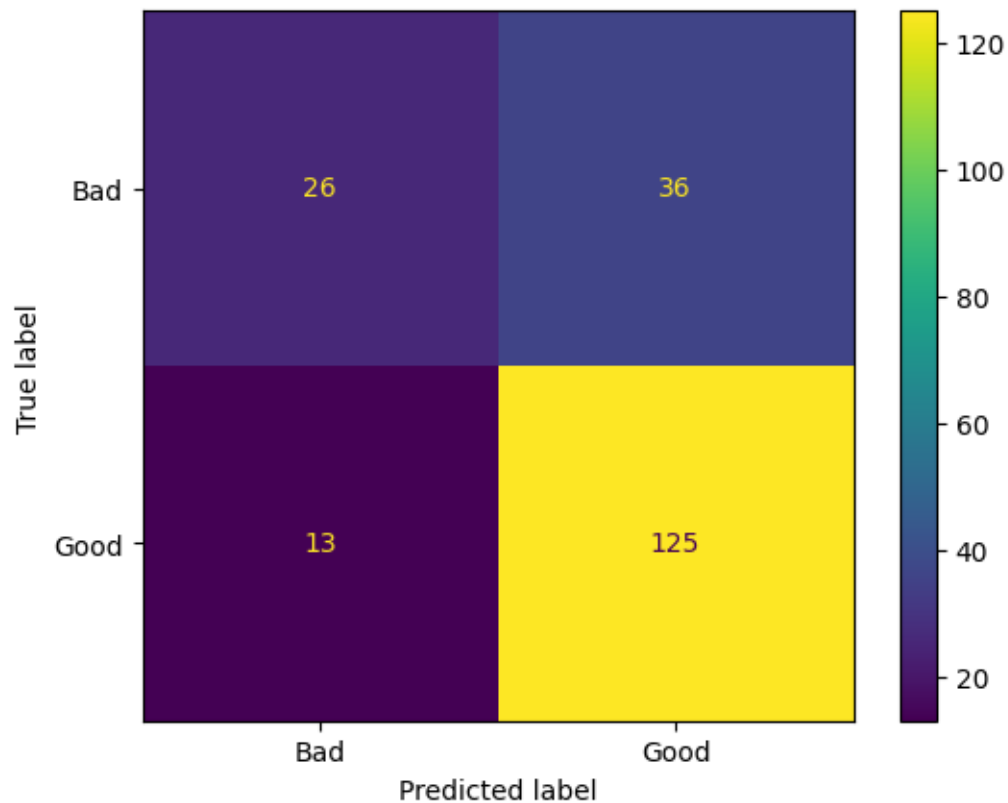
preds_new_features = rf.predict(selected_features_test)
cr_feature = classification_report(y_test, preds_new_features)
print(cr_feature)

ConfusionMatrixDisplay.from_predictions(y_test, preds_new_features)
```

```
precision    recall  f1-score   support
```

Bad	0.67	0.42	0.51	62
Good	0.78	0.91	0.84	138
accuracy			0.76	200
macro avg	0.72	0.66	0.68	200
weighted avg	0.74	0.76	0.74	200

[61]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x318a90950>



```
[62]: # Hypertuning parameters
param_grid = {
    'n_estimators': list(np.arange(25, 200, 25)),
    'max_features': ['sqrt', 'log2', None],
    'max_depth': list(np.arange(3, 21, 3)),
    'max_leaf_nodes': list(np.arange(3, 21, 3)),
}

search = RandomizedSearchCV(RandomForestClassifier(), param_grid)
search.fit(selected_features_train, y_train)
```

```

print(search.best_estimator_)

# Retraining with tuned hyperparameters
rf = RandomForestClassifier(max_depth=6, max_features='log2',
    ↪max_leaf_nodes=18, n_estimators=75)
rf.fit(selected_features_train, y_train_flat)

# Predictions with tuned hyperparameters

preds_new_features_tuned = rf.predict(selected_features_test)
cr_feature_tuned = classification_report(y_test, preds_new_features_tuned)
print(cr_feature_tuned)

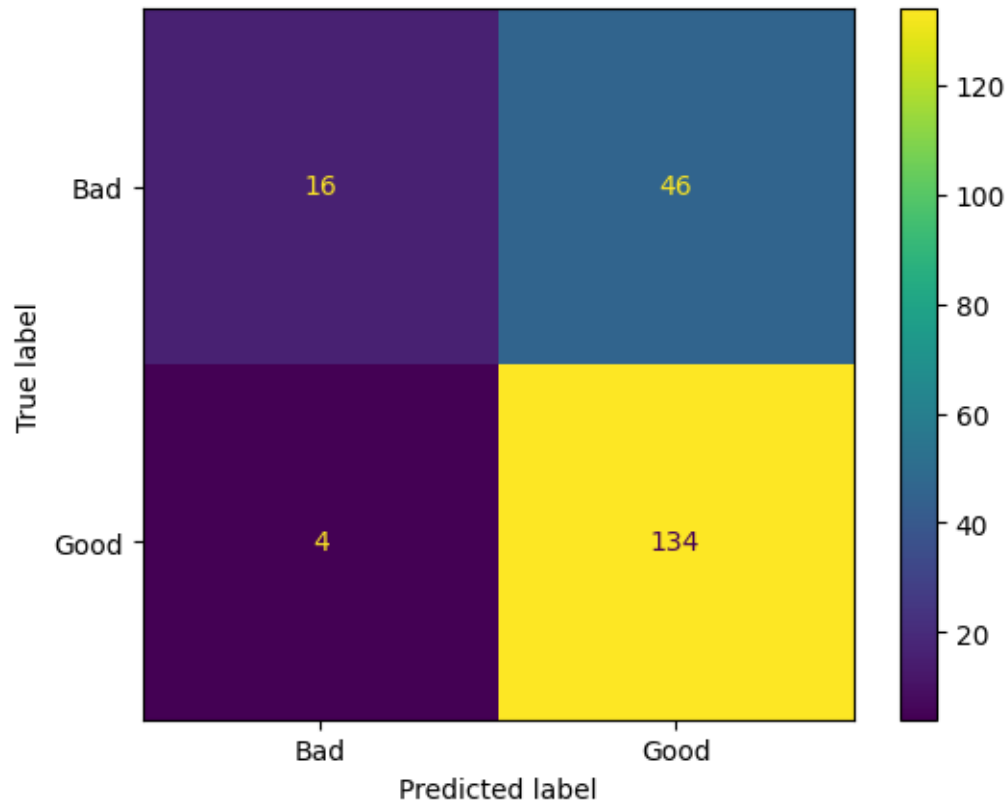
ConfusionMatrixDisplay.from_predictions(y_test, preds_new_features_tuned)

```

```
RandomForestClassifier(max_depth=6, max_leaf_nodes=18, n_estimators=50)
```

	precision	recall	f1-score	support
Bad	0.80	0.26	0.39	62
Good	0.74	0.97	0.84	138
accuracy			0.75	200
macro avg	0.77	0.61	0.62	200
weighted avg	0.76	0.75	0.70	200

[62]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x31897da50>



0.8.1 Random Forest Interpretation

The RandomForest model demonstrates a strong ability to correctly identify positive predictions (Precision) and to capture a significant portion of the actual positive cases (Recall), both with equal proficiency. The model achieves a balanced performance as evidenced by its F1 score, which is commendable though slightly lower than the individual precision and recall scores. The overall accuracy, which measures the model's effectiveness in identifying both True Positives and True Negatives, aligns closely with the other metrics, highlighting its consistency across various aspects of classification. The results indicate a robust and reliable model performance.

0.8.2 Data Preparation Impact on Model Performance

RandomForest overall performance landed it in the middle of Naive Bayes and GBM. This is expected since in general, that is the order of "Performance power" of each given model. And since Random Forest still reached second place despite less data cleaning, it brings into question the overall value of the initial data cleaning. One possible reason for the lack of major difference could be the overall quality of the data and had the data been "dirtier" then perhaps a major difference would have been observed.

0.9 Conclusion

In this analysis, we compared the performance of three machine learning models in predicting good or bad credit applicants: Naive Bayes, RandomForest, and Gradient Boosting Machine (GBM) across various metrics. GBM demonstrated the strongest performance, particularly in precision, recall, and F1 score, suggesting it is highly effective for complex prediction tasks like assessing creditworthiness. RandomForest (with less data cleaning) showed robustness and consistency, making it a reliable choice, especially for datasets with complex features. Naive Bayes, while fastest and effective in smaller datasets, lagged slightly behind in performance due to its assumption of feature independence. Given this conclusion, we have decided that the Gradient Boosting Machine (GBM) model would be the best model to move forward with in predicting good or bad credit applicants.